



# Introduction

ITEV DBMS – 1<sup>st</sup> Seminar, Sep. 6, 2008

**Original slides are from**  
**Silberschatz, Korth and Sudarshan**





# Today's Agenda

- 09.00 – 09.30 Introduction
- 09.30 – 11.00 Relational Model
- 11.00 – 12.00 SQL (1)
  
- 12.00 – 12.45 Lunch
  
- 12.45 – 13.45 SQL (2)
- 13.45 – 15.00 Entity-Relationship Model
- 15.00 – 15.45 Mini-Project, exam and assignments
- 15.45 – 16.00 Summary of the day





# Administration Info

- Course lecturer
  - Hua Lu ([luhua@cs.aau.dk](mailto:luhua@cs.aau.dk), Tel: 9940 9973)
- Course website
  - <http://www.cs.aau.dk/~luhua/courses/itev-db08/>
- Seminar dates
  - September 6, 2008
  - September 13, 2008
  - October 4, 2008
- Oral exam and mini-project
  - Exam date: October 18, 2008
  - Mini-project report deadline: October 16, 2008 at 14.00
    - ▶ 2-4 people a group is preferred
    - ▶ Form groups in the lunch break if possible (deadline: Monday, Sept. 8)
    - ▶ **No report, no exam!**
  - More details about exam and mini-project in the afternoon






# Course Structure

- Introduction to DBMS
- Relational Databases and Database Design

## Day 1

- Relational Model
- SQL
- Entity-Relationship (E-R) Model
- Normalization



- Data Storage and Querying

## Day 2

- Storage and File Structure
- Indexing and Hashing
- Query Processing and Optimization



- Transaction Management

## Day 3

- Transactions
- Concurrency Control
- Recovery





# Introduction to DBMS

- Purpose of Database Systems
- View of Data
- Database Languages
- Relational Databases
- Database Design
- Data Storage and Querying
- Transaction Management





# Database Management System (DBMS)

- DBMS contains information about a particular organization
  - Collection of interrelated data
  - Set of programs to access the data
  - An environment that is both *convenient* and *efficient* to use
- Database Applications:
  - Banking: all transactions
  - Airlines: reservations, schedules
  - Universities: registration, grades
  - Sales: customers, products, purchases
  - Online retailers: order tracking, customized recommendations
  - Manufacturing: production, inventory, orders, supply chain
  - Human resources: employee records, salaries, tax deductions
- Databases touch all aspects of our lives





# Purpose of Database Systems

- In the early days, database applications were built directly on top of file systems
- Drawbacks of using file systems to store data:
  - Data redundancy and inconsistency
    - ▶ Multiple file formats, duplication of information in different files
  - Difficulty in accessing data
    - ▶ Need to write a new program to carry out each new task
  - Data isolation — multiple files and formats
    - ▶ Interoperability issue
  - Integrity problems
    - ▶ Integrity constraints (e.g. account balance  $> 0$ ) become “buried” in program code rather than being stated explicitly
    - ▶ Hard to add new constraints or change existing ones





# Purpose of Database Systems (Cont.)

- Drawbacks of using file systems (cont.)
  - Atomicity of updates
    - ▶ Failures may leave database in an inconsistent state with partial updates carried out
    - ▶ Example: Transfer of funds from one account to another should either complete or not happen at all
  - Concurrent access by multiple users
    - ▶ Concurrent accessed needed for performance
    - ▶ Uncontrolled concurrent accesses can lead to inconsistencies
      - Example: Two people reading a balance and updating it at the same time
  - Security problems
    - ▶ Hard to provide user access to some, but not all, data
- Database systems offer solutions to all the above problems







# Levels of Abstraction

- **Physical level:** describes how a record (e.g., customer) is stored.
- **Logical level:** describes data stored in database, and the relationships among the data.

**type** *customer* = **record**

```
customer_id : string;  
customer_name : string;  
customer_street : string;  
customer_city : integer;
```

**end;**

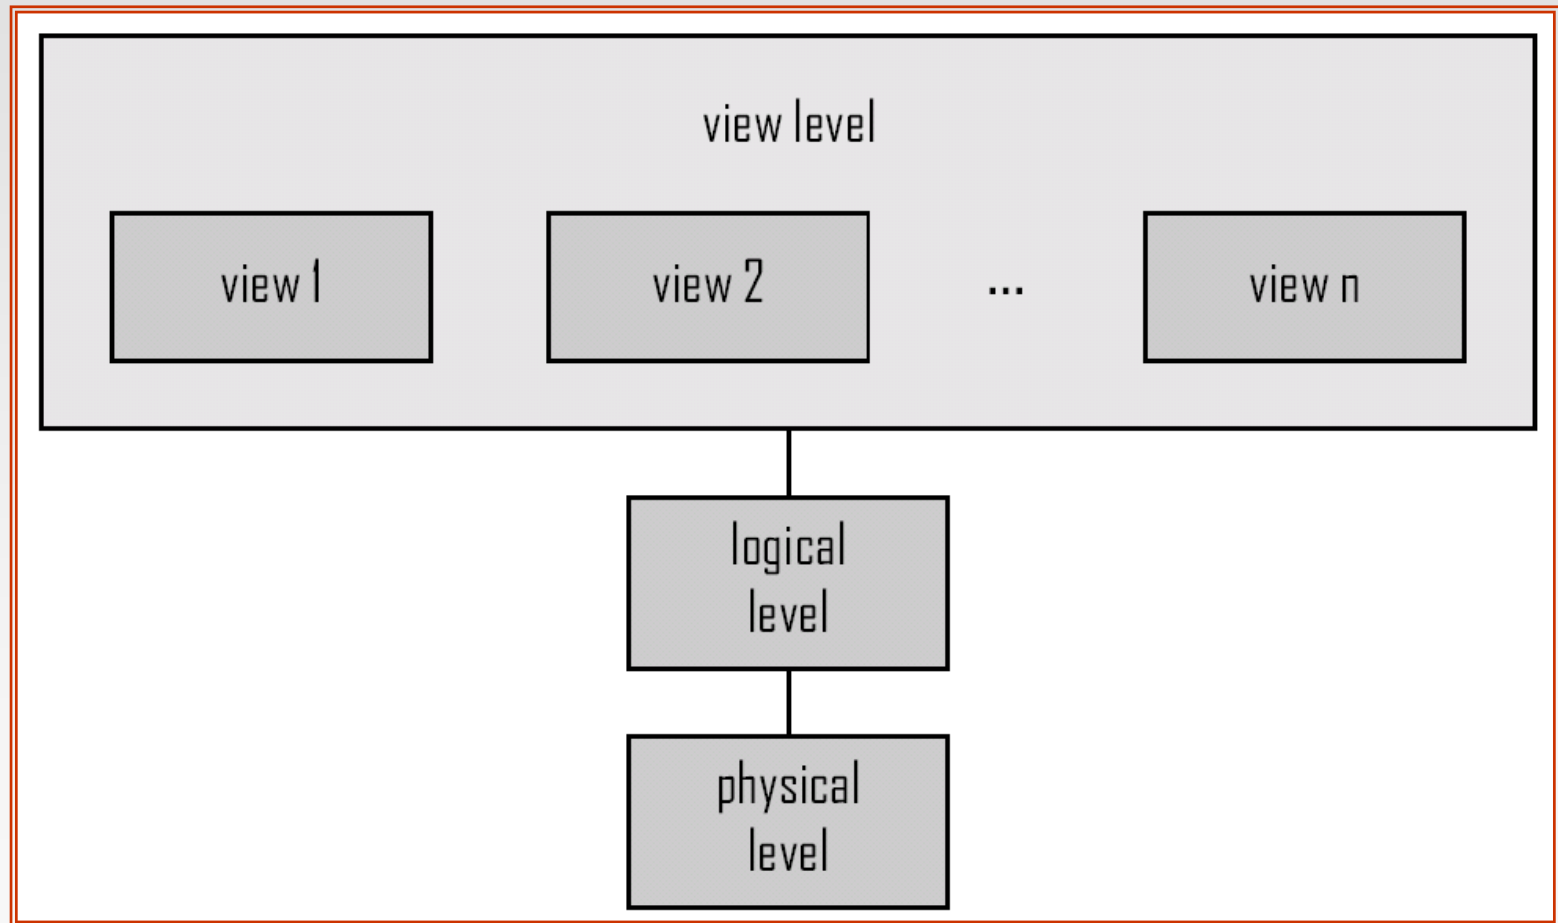
- **View level:** application programs hide details of data types. Views can also hide information (such as an employee's salary) for security purposes.





# View of Data

An architecture for a database system





# Instances and Schemas

- Similar to types and variables in programming languages
- **Schema** – the logical structure of the database
  - Example: The database consists of information about a set of customers and accounts and the relationship between them)
  - Analogous to type information of a variable in a program
  - **Physical schema**: database design at the physical level
  - **Logical schema**: database design at the logical level
- **Instance** – the actual content of the database at a particular point in time
  - Analogous to the value of a variable
- **Physical Data Independence** – the ability to modify the physical schema without changing the logical schema
  - Applications depend on the logical schema
  - In general, the interfaces between the various levels and components should be well defined so that changes in some parts do not seriously influence others.





# Data Models

- A collection of tools for describing
  - Data
  - Data relationships
  - Data semantics
  - Data constraints
- Relational model
- Entity-Relationship data model (mainly for database design)
- Object-based data models (Object-oriented and Object-relational)
- Semistructured data model (XML)
- Other older models:
  - Network model
  - Hierarchical model





# Data Definition Language (DDL)

- Specification notation for defining the database schema

Example:     **create table** *account* (  
                            *account-number*   **char**(10),  
                            *balance*           **integer**)

- DDL compiler generates a set of tables stored in a *data dictionary*
- Data dictionary contains metadata (i.e., data about data)
  - Database schema
  - Data *storage and definition* language
    - ▶ Specifies the storage structure and access methods used
  - Integrity constraints
    - ▶ Domain constraints
    - ▶ Referential integrity (**references** constraint in SQL)
    - ▶ Assertions
  - Authorization





# Data Manipulation Language (DML)

- Language for accessing and manipulating the data organized by the appropriate data model
  - DML also known as query language
- Two classes of DMLs
  - **Procedural DML** – user specifies what data is required and how to get those data
  - **Declarative (nonprocedural) DML** – user specifies what data is required without specifying how to get those data
- SQL is the most widely used query language
  - Set-based, declarative
  - But procedural extensions are offered by different database systems





# Relational Model

- Example of tabular data in the relational model

Attributes

<i>customer_id</i>	<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>	<i>account_number</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto	A-101
192-83-7465	Johnson	12 Alma St.	Palo Alto	A-201
677-89-9011	Hayes	3 Main St.	Harrison	A-102
182-73-6091	Turner	123 Putnam St.	Stamford	A-305
321-12-3123	Jones	100 Main St.	Harrison	A-217
336-66-9999	Lindsay	175 Park Ave.	Pittsfield	A-222
019-28-3746	Smith	72 North St.	Rye	A-201





# A Sample Relational Database

<i>customer_id</i>	<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto
677-89-9011	Hayes	3 Main St.	Harrison
182-73-6091	Turner	123 Putnam Ave.	Stamford
321-12-3123	Jones	100 Main St.	Harrison
336-66-9999	Lindsay	175 Park Ave.	Pittsfield
019-28-3746	Smith	72 North St.	Rye

(a) The *customer* table

<i>account_number</i>	<i>balance</i>
A-101	500
A-215	700
A-102	400
A-305	350
A-201	900
A-217	750
A-222	700

(b) The *account* table

<i>customer_id</i>	<i>account_number</i>
192-83-7465	A-101
192-83-7465	A-201
019-28-3746	A-215
677-89-9011	A-102
182-73-6091	A-305
321-12-3123	A-217
336-66-9999	A-222
019-28-3746	A-201

(c) The *depositor* table







# SQL

- **SQL**: widely used non-procedural language
  - Example: Find the name of the customer with customer-id 192-83-7465

```
select   customer.customer_name
from     customer
where    customer.customer_id = '192-83-7465'
```
  - Example: Find the balances of all accounts held by the customer with customer-id 192-83-7465

```
select   account.balance
from     depositor, account
where    depositor.customer_id = '192-83-7465' and
          depositor.account_number = account.account_number
```
- Application programs generally access databases through one of
  - Language extensions to allow embedded SQL
  - Application program interface (e.g., ODBC/JDBC) which allow SQL queries to be sent to a database





# Database Design

The process of designing the general structure of the database:

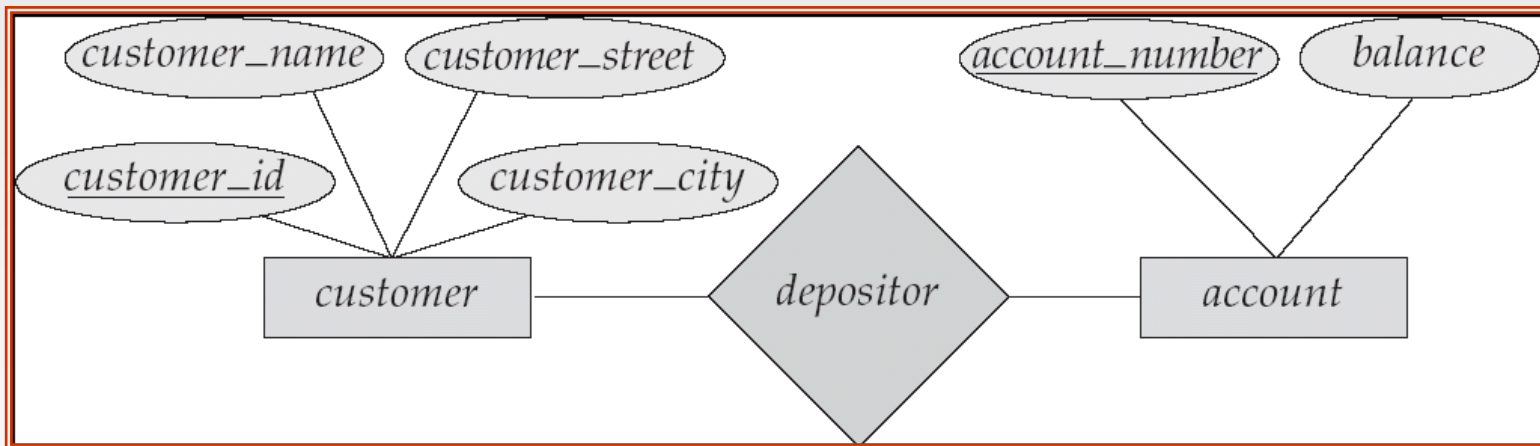
- **Conceptual and Logical Designs** – Deciding on the database schema. Database design requires that we find a “good” collection of relation schemas.
  - **Business decision** – What attributes should we record in the database?
  - **Computer Science decision** – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
- **Physical Design** – Deciding on the physical layout of the database





# The Entity-Relationship Model

- Models an organization as a collection of *entities* and *relationships*
  - Entity: a “thing” or an “object” in the organization that is distinguishable from other objects
    - ▶ Described by a set of *attributes*
  - Relationship: an association among several entities
- Represented diagrammatically by an *entity-relationship diagram*:





# Storage Management

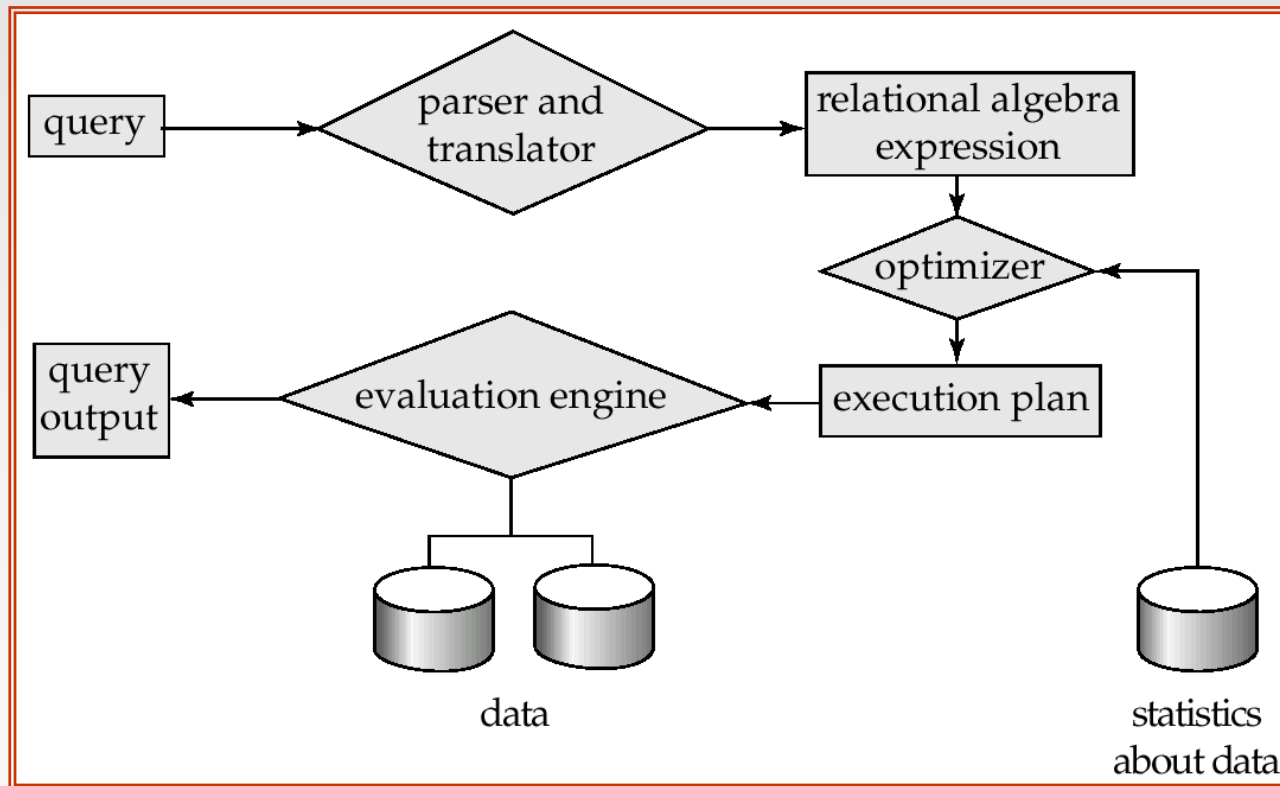
- **Storage manager** is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.
- The storage manager is responsible to the following tasks:
  - Interaction with the file manager
  - Efficient storing, retrieving and updating of data
- Issues:
  - Storage access
  - File organization
  - Indexing and hashing





# Query Processing

1. Parsing and translation
2. Optimization
3. Evaluation





# Query Processing (Cont.)

- Alternative ways of evaluating a given query
  - Equivalent expressions
  - Different algorithms for each operation
- Cost difference between a good and a bad way of evaluating a query can be enormous
- Need to estimate the cost of operations
  - Depends critically on statistical information about relations which the database must maintain
  - Need to estimate statistics for intermediate results to compute cost of complex expressions





# Transaction Management

- A **transaction** is *a collection of operations* that performs a single logical function in a database application
- **Transaction management component** ensures that the database remains in a consistent (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.
- **Concurrency control manager** controls the interaction among the concurrent transactions, to ensure the consistency of the database.





# End of Introduction

ITEV DBMS – 1<sup>st</sup> Seminar, Sep. 6, 2008

## Acknowledgements

- Slides from the SKS book site
- Slides from Kristian Torp, course lecturer in 2006

**Original slides are from**  
**Silberschatz, Korth and Sudarshan**

