

Concurrency Control: Outline

- Two-Phase Locking
 - Basics and Motivation
 - Lock conversion
 - Strict two-phase protocol
- Graph-Based Locking
- Deadlock Handling
 - Deadlock Detection
 - Deadlock Prevention
- Timestamp-Based Protocol
- Lock Tuning

Concurrency Control

- Two-Phase Locking
 - Basics and Motivation
 - Lock conversion
 - Strict two-phase protocol
- Graph-Based Locking
 - Simple Tree Protocol
- Deadlock Handling
- Timestamp-Based Protocol
- Lock Tuning

A Locking Protocol

- This protocol ensures only serializable schedules by delaying transactions that could violate serializability.
- Two types of *locks* can be set on an item Q:
 - *Exclusive lock*: LX Q
 - *Shared lock*: LS Q
- Locks can be released.
 - *Unlock*: UN Q
- Privileges associated with locks
 - A transaction holding an X-lock may issue a write or read access request on the item.
 - A transaction holding as S-lock may issue a read access request on the item.

Compatibility Matrix

- This matrix is interpreted as follows: a transaction may set a lock on an item if this lock is compatible with locks already held on the item by other transactions.

| | | <i>New</i> | |
|------------|---|------------|-------|
| | | X | S |
| <i>Old</i> | X | False | False |
| | S | False | True |

- It follows that
 - Any number of transactions can hold S-locks on an item.
 - If any transaction holds an X-lock on the item, no other transaction may hold any lock on the item.

Lock Granting

- Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.
- A transaction may be granted a lock on an item if the requested lock is *compatible* with locks already held on the item by other transactions
- If a lock cannot be granted, the requesting transaction is made to wait till all *incompatible* locks held by other transactions have been released. The lock is then granted.

Problems with Locking

- Transaction *T15* transfers \$50 from account B to account A.
- Transaction *T16* displays the total amount of money in accounts A and B.

T15:

| |
|-----------------------|
| LX B |
| read B |
| $B \leftarrow B - 50$ |
| write B |
| UN B |
| LX A |
| read A |
| $A \leftarrow A + 50$ |
| write A |
| UN A |

T16:

| |
|-------------|
| LS A |
| read A |
| UN A |
| LS B |
| read B |
| UN B |
| display A+B |

- Initially $A = 100$ and $B = 200$.

Problems with Early Unlocking

| Schedule S7 | |
|--|---|
| T15 | T16 |
| LX B read B $B \leftarrow B - 50$ write B UN B | |
| | LS A read A UN A LS B read B UN B display A + B |
| LX A read A $A \leftarrow A + 50$ write A UN A | |

- Early unlocking can cause incorrect results (non-serializable schedules).

Problems with Late Unlocking

- So let's delay unlocking until the end of the transaction.

| Schedule S8 | |
|---------------------------------------|------------------------|
| T17 | T18 |
| LX B read B B ← B-50 write B | |
| | LS A read A LS B |
| LX A | |

- Late unlocking can cause *deadlocks*.

The Two-Phase Locking Protocol

- First phase (*growing phase*):
 - Transaction may request locks.
 - Transaction may not release locks.
- Second phase (*shrinking phase*):
 - Transaction may not request locks.
 - Transaction may release locks.
- When the first lock is released, the transaction moves from the first phase to the second phase.

Precedes Relation for Locks

- A locking protocol ensures serializability if and only if for all legal schedules under the protocol, the *precedes relation* is acyclic.
- T_i precedes T_j ($T_i \rightarrow T_j$) for schedule S if
 - i does not equal j .
 - T_i read/write locks Q , T_j is the next transaction to write lock Q .
 - T_i write locks Q . T_j is the next transaction that read locks Q (after T_i unlocks Q but before any other transaction write locks Q).
- To discover that a (locking) schedule is serializable
 - Construct a conflict graph using the definition above.
 - Check acyclicity.
- Topological sort gives a serialization order.

Two-Phase Locking Works

- **Theorem** If S is any schedule of two-phase transactions, then S is serializable.
- **Proof**
 - Suppose not. Then conflict graph must have a cycle.
$$T_i \rightarrow T_j \rightarrow \dots \rightarrow T_m \rightarrow T_i$$
 - A lock by T_j follows an unlock by T_i , and a lock by T_i follows an unlock by T_m . By transitivity, a lock by T_i follows an unlock by T_i .
 - Contradiction.
 - ◆ Not consistent with the two-phase locking protocol definition.

Generality of Two-Phase Locking

- In general, a locking protocol does not allow all serializable schedules. Some serializable schedules are illegal.
- *S9* is an illegal schedule under the two-phase protocol.

| Schedule <i>S9</i> | | |
|--------------------|--------------|--------------|
| <i>T19</i> | <i>T20</i> | <i>T21</i> |
| LX B UN B | LX A UN A | LX A UN A |
| | LX B UN B | |

- Serialization order: $T19 \rightarrow T20 \rightarrow T21$

Two-Phase With Lock Conversion

- First Phase:
 - Can acquire an S-lock on item.
 - Can acquire an X-lock on item.
 - Can *convert (upgrade)* an S-lock to an X-lock.
- Second Phase:
 - Can release an S-lock.
 - Can release an X-lock.
 - Can *convert (downgrade)* an X-lock to an S-lock.
- This protocol assures serializability.
- It relies on the application programmer to insert the appropriate locks.

Extended Conflict Graph

- T_i precedes T_j ($T_i \rightarrow T_j$) if there exists a data-item Q such that
 - T_i has held lock-mode A on Q .
 - T_j has held lock-mode B on Q .
 - T_j 's lock is later than T_i 's lock.
 - $\text{COMP}(A, B) = \text{false}$. (COMP is the compatibility matrix.)
- T_i must occur before T_j in a serial schedule.

Practical Protocol

- The programmer (transaction) issues the standard read/write instructions.
- The system manages all aspects of the protocol, including the lock operations.

-

read D:

if T has a lock on D
 perform read

else

 wait until no other transaction has a X-lock on D
 grant T an S-lock on D
 perform read

Practical Protocol, cont.

write D:

if T has a X-lock on D

perform write

else

wait until no other transaction has a lock on D

if T has an S-lock on D

convert it to X-lock

else

grant it X-lock on D

perform write

- All locks are released after commit or abort (Rollback)

Cascading Rollbacks

- One transaction aborting can cause other transactions to abort.

| Schedule S11 | | |
|----------------------|--------------|------|
| T22 | T23 | T24 |
| LX A LX B UN A | LX A UN A | LX A |

- $T22$ aborts \Rightarrow we have to rollback $T23$ and $T24$.
- How to eliminate these *cascading rollbacks*?
 - Don't let transactions read “dirty” uncommitted data.

Strict Two-Phase Locking

- In **strict two-phase locking**, *x-locks* are not released until after commit time.
- In **rigorous two-phase locking**, *all locks* are release after commit time.
- Hence, no cascading rollbacks.
- Loss of potential concurrency.

Concurrency Control

- Two-Phase Locking
- Graph-Based Locking
 - Simple Tree Protocol
- Deadlock Handling
- Timestamp-Based Protocol
- Lock Tuning


Graph-Based Protocols

- Graph-based protocols are an alternative to two-phase locking.
- Impose a partial ordering \rightarrow on the set $\mathbf{D} = \{d_1, d_2, \dots, d_h\}$ of all data items.
 - If $d_i \rightarrow d_j$ then any transaction accessing both d_i and d_j must access d_i before accessing d_j .
 - Implies that the set \mathbf{D} may now be viewed as a directed acyclic graph, called a *database graph*.
- The *tree-protocol* is a simple kind of graph protocol.

Simple Tree Protocol

- Simple Tree Protocol

- Partial order on DB items determines a tree.
- Only exclusive locks are allowed.

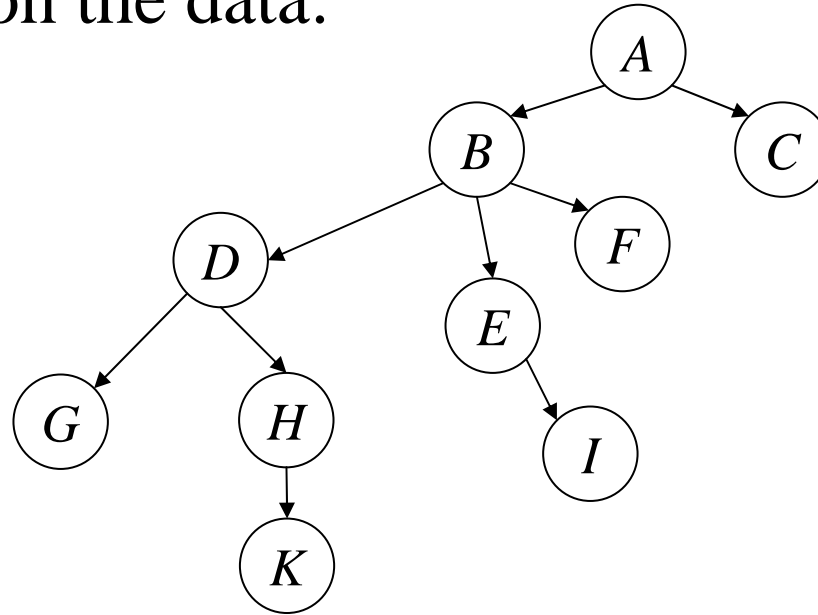


- Rules

- Each T_i can lock a data item at most once.
- First lock on any data item.
- Subsequently, T_i can lock item Q only if it currently holds a lock on the parent of Q.
- Data items can be unlocked any time.


Simple Tree Protocol Example

- Partial order on the data:



- LX B; LX E; UN E; LX D; UN B; LX G; UN D; UN G.**
- LX D; LX H; UN D; LX K; UN K; UN H.**
- LX B; LX E; UN E; UN B.**
- LX D; LX H; UN D; UN H.**

Simple Tree Protocol Example, cont.



| Schedule S12 | | | |
|------------------------------|----------------------|--------------|------------------------------|
| T25 | T26 | T27 | T28 |
| LX B | LX D LX H UN D | | |
| LX E UN E LX D UN B | | LX B LX E | |
| LX G UN D | LX K UN K UN H | | |
| UN G | | UN E UN B | LX D LX H UN D UN H |

Concurrency Control

- Two-Phase Locking
- Graph-Based Locking
- Deadlock Handling
 - Deadlock Detection
 - Deadlock Prevention
- Timestamp-Based Protocol
- Lock Tuning


Deadlocks

- “Easy” solutions:
 - Newer wait and do (partial) rollback, may lead to *livelock*.
 - ◆ The states of the transactions involved constantly change with regard to one another, none progressing.
 - Linearly order all resources (preventions technique).
- Alternatives for deadlock handling
 - Detect deadlock and resolve
 - Prevent
 - Timeout
 - ◆ All systems must ultimately depend on timeout to detect some deadlocks.

Deadlock Detection

- Create a “wait for” graph and check for cycles.
 - An edge from $T1$ to $T2$ if $T1$ waits for $T2$
- Invoke detection algorithm repeatedly.
- If deadlock is detected,
 - Select appropriate victim.
 - Abort victim and release its locks.

| Schedule S8 | |
|--|------------------------|
| $T17$ | $T18$ |
| LX B read B $B \leftarrow B-50$ write B | |
| | LS A read A LS B |
| LX A | |



Deadlock Detection Example

| <i>T1</i> | <i>T2</i> | <i>T3</i> | <i>T4</i> |
|--------------|-----------|-----------|-----------|
| | LX A | | |
| LX B LX C | | | |
| LX D | | LX B | |
| | LX D | | LX E |
| LX E | | | LX A |

Exercise:

Draw out the “wait for” graph, and detect if deadlock exists.

Hint:

List in a table all resources, locks on each, and requests on each.

Deadlock Prevention Protocols

- Give each transaction a timestamp (separate from any timestamp used for concurrency control) that is retained even across rollbacks.
- Assume T_i requests an item held by T_j .
- *Wait-Die*
 - Wait: If T_i is older than T_j , then T_i waits.
 - Die: If T_i is younger than T_j , then abort T_i .
- *Wound-Wait*
 - Wound: If T_i is older than T_j , then abort T_j .
 - Wait: If T_i is younger than T_j , then T_i waits.
- No starvation in either scheme, because transactions advance.

Wait-Die Examples

T10 is older than T20.

| <i>T10</i> | <i>T20</i> |
|------------|------------|
| LX A | LX A |
| | ROLLBACK |
| | LX A |
| | ROLLBACK |
| | LX A |
| | ROLLBACK |
| UN A | LX A |
| | ... |
| | UN A |

T10 starts first.

| <i>T10</i> | <i>T20</i> |
|------------|------------|
| | LX A |
| LX A | |
| WAIT | |
| WAIT | |
| WAIT | |
| WAIT | |
| WAIT | |
| | UN A |
| LX A | |
| ... | |
| UN A | |

T20 starts first.

Wound-Wait Examples

T10 is older than T20.

| <i>T10</i> | <i>T20</i> |
|------------|------------|
| LX A | |
| | LX A |
| | WAIT |
| | WAIT |
| | WAIT |
| | WAIT |
| | WAIT |
| UN A | |
| | LX A |
| | ... |
| | UN A |

T10 starts first.

| <i>T10</i> | <i>T20</i> |
|------------|------------|
| | LX A |
| LX A | |
| | ROLLBACK |
| | LX A |
| | WAIT |
| | WAIT |
| | WAIT |
| UN A | |
| | LX A |
| | ... |
| | UN A |

T20 starts first.

Concurrency Control

- Two-Phase Locking
- Graph-Based Locking
- Deadlock Handling
- Timestamp-Based Protocol
- Lock Tuning

Timestamp Protocol

- Each transaction is issued a timestamp when it enters the system.
- Timestamps are drawn from an increasing sequence of integers.
- The protocol will manage the concurrent execution so that it will be equivalent to a serial execution in timestamp order.
- If two transactions conflict in the schedule, then the protocol will ensure that the one with the lower timestamp accesses the item first.

The Timestamp Protocol, cont.

- In order to assure such behavior the protocol maintains for each item Q two integers, used for synchronization.
 - $r-max(Q)$: the maximal timestamp of a transaction that read that item.
 - $w-max(Q)$: the maximal timestamp of a transaction that wrote the item.
- If a transaction wants to access an item in a way that will not create a prohibited conflict, the access is allowed and a synchronization value is updated.
- If the conflict is prohibited, the transaction is aborted, and restarted with a *new (later) timestamp*.

Implementing the Timestamp Protocol

- Let t be the timestamp of the executing transaction.

read X: **if** $t \geq w\text{-max}(X)$
 perform read
 $r\text{-max}(X) \leftarrow \max(t, r\text{-max}(X))$
 else abort

write X: **if** $t \geq r\text{-max}(X)$ **and** $t \geq w\text{-max}(X)$
 perform write
 $w\text{-max}(X) \leftarrow t$
 else
 if $t < r\text{-max}(X)$
 abort
 else
 do nothing

Timestamp Protocol Example

- We present a partial history for several items for transactions with timestamps 1, 2, 3, 4, and 5.

| Schedule S13 | | | | |
|--------------|-----------------|--------------------|------------------|--------------------|
| T29 | T30 | T31 | T32 | T33 |
| read Y | read Y | write Y write Z | | read X |
| read X | read Z abort | | | read Z |
| | | | write Z abort | write X write Z |

Exercise:

Figure out $r\text{-max}(Q)$ and $w\text{-max}(Q)$ values after each step, for $Q=X, Y, Z$

Timestamp vs. Locking

- Schedule allowed by locks but not timestamps.

| <i>Schedule S14</i> | |
|---------------------|-------------------|
| <i>T34</i> | <i>T35</i> |
| read A | read B write B |
| read B | |

- Schedule allowed by timestamps but not by locks.

| <i>Schedule S15</i> | | |
|---------------------|------------------------|------------|
| <i>T36</i> | <i>T37</i> | <i>T38</i> |
| write A | write A write B | write A |
| write B | | |

Strict Timestamp Based Concurrency Control

- How to avoid cascading rollbacks?
 - Transactions should read only committed values.
- *Strict timestamp concurrency control protocol*
- *read X:*
 - if** $t \geq w\text{-max}(X)$
 - $r\text{-max}(X) \leftarrow \max(t, r\text{-max}(X))$
 - wait for a committed value of X
 - perform read
 - else** abort
- *write X:*
 - if** $t \geq r\text{-max}(X)$ **and** $t \geq w\text{-max}(X)$
 - $w\text{-max}(X) \leftarrow t$
 - wait until X value is a committed value and pending reads are done
 - perform write
 - else**
 - if** $t < r\text{-max}(X)$ abort
 - else** do nothing

Concurrency Control

- Two-Phase Locking
- Graph-Based Locking
- Deadlock Handling
- Timestamp-Based Protocol
- Lock Tuning

Lock Tuning: Various Techniques

- Eliminate locking when it is unnecessary.
 - Switch off locking when only one transaction is running, e.g., when loading a database.
 - Eliminate locking when all transactions are read-only (e.g., in Data Warehousing)
- Apply transaction chopping/reduce the length of transactions.
 - A long transaction must wait for locks and makes other transactions wait for locks.
 - Interactive transactions that involve human interaction, e.g., airline reservations, may be broken into separate transactions.
 - Chopping may affect correctness if one is not careful.
- Avoid DDL statements.
 - DDL statements update the system catalog, which all transactions must access.
 - DDL statements thus tend to slow down transactions.

Use Different Isolation Levels

- Consider weakening isolation (correctness) when the application permits it.
- Isolation levels:
 - **Degree 0:** Reads may see uncommitted ("dirty") data (yielding "dirty reads"). Reads are *non-repeatable*. Overwriting of dirty data is possible.
 - **Degree 1:** Dirty and non-repeatable reads, but no overwriting of dirty data.
 - **Degree 2:** Non-repeatable reads allowed, but no dirty reads. Using standard locking, read locks are released immediately; write locks are managed according to the two-phase protocol.
 - **Degree 3:** Reads are now also repeatable. ACID transactions result.
- Guideline:
 - Degree 3 is the default. If transactions are long and there is a lot of deadlocks and blocking, consider using Degree 2 or 1.

Select the Right Lock Granularity

- Three levels of locks are common:
 - record-level, page-level, and table-level locks
- Rule of thumb
 - If every transaction accesses well below 1% of the records of a table, and the records are on different pages, record-level locking is best.
 - ◆ This is the case in environments with many short transactions.
 - Medium-length transactions that use clustering indices should use page-level locking.
 - Long transactions should use table-level locking.

Hotspots

- Avoid hotspots.
 - Monitoring tools are used for locating hotspots.
 - Use partitioning.
 - Access hotspots late in a transaction.
- Partitioning example:
 - Suppose sequential keys create concurrency control bottlenecks.
 - ◆ Use many insertion points and insert randomly.
 - ◆ Cluster on an uncorrelated attribute.

Summary

- Many concurrency control protocols have been developed.
- There are several entire books on the topic.
- Most relational DBMS's use rigorous two-phase locking.
- Many refinements, such as page and record locking, are used.
- Tuning is important for good performance.