

Recovery: Outline

- Logging Mechanism
 - Motivation
 - Log-Based Recovery
- Recovery Algorithms
 - Undo/Redo
 - No-Undo/Redo
 - Undo/No-Redo
 - No-Undo/No-Redo
- Checkpoint
- Other Issues
 - Media Failure and Archiving
 - Recovery Tuning


Recovery

- Logging Mechanism
 - Motivation
 - Log-Based Recovery
- Recovery Algorithms
- Checkpoint
- Other Issues

Motivation

- Atomicity:
 - Transactions may abort (“Rollback”).
- Durability:
 - What if DBMS stops running? (Causes?)

- ❖ Desired Behavior after system restarts:
 - T1, T2 & T3 should be durable.
 - T4 & T5 should be aborted (effects not seen).



Atomicity and Durability of Operations

- Each database operation
 - either completes fully,
 - or does not happen at all.
- Low-level atomic operations are built into hardware.
- High-level atomic operations are called *transactions*.
- The DBMS ensures that a transaction
 - either completes and has a permanent result (*committed* transaction),
 - or has no effect at all on the database (*aborted* transaction).
- The role of recovery is to ensure atomicity and durability of transactions in the presence of system failures.

A Sample Transaction

- Transfer \$50 from account A to account B.
 - 1 Transaction starts
 - 2 Read A
 - 3 $A \leftarrow A - 50$
 - 4 Write A
 - 5 Read B
 - 6 $B \leftarrow B + 50$
 - 7 Write B
 - 8 Transaction ends
- If initial values are $A = 180$, $B = 100$, then after the execution $A = 130$ and $B = 150$.
- Sample values of database variables at various points in execution:

Last Instruction	A	B
2	180	100
5	130	100
7	130	150



Problems In Ensuring Atomicity and Durability

- *Logical errors*: A transaction cannot complete due to some internal condition, e.g., bad input, overflow, resource limit exceeded.
- *System errors*: The database system must terminate an active transaction due to, e.g., deadlock. The transaction can be executed later.
- *System crash*: A power failure or other hardware failure causes the system to crash. Volatile storage is lost, non-volatile is not.
- *Disk failure*: A head crash or similar failure destroys all or part of disk storage (failure during data transfer, fire, theft, sabotage, mounting wrong disk, etc.). Nonvolatile storage is lost.

Recovery

- Initial (incorrect) procedures
 - Re-execute the transaction.
 - ◆ If the sample transaction crashed after instruction 4 (or later), incorrect values ($A=130, \dots$) will exist in the database.
 - Do not re-execute the transaction.
 - ◆ If the sample transaction crashed before instruction 7, incorrect values ($\dots, B = 100$) will exist in the database.
- Problems
 - Which instruction was last executed?
 - Multiple, concurrent users
 - Buffer management

A = 180, B = 100

- 1 Transaction starts
- 2 Read A
- 3 $A \leftarrow A - 50$
- 4 Write A
- 5 Read B
- 6 $B \leftarrow B + 50$
- 7 Write B
- 8 Transaction ends

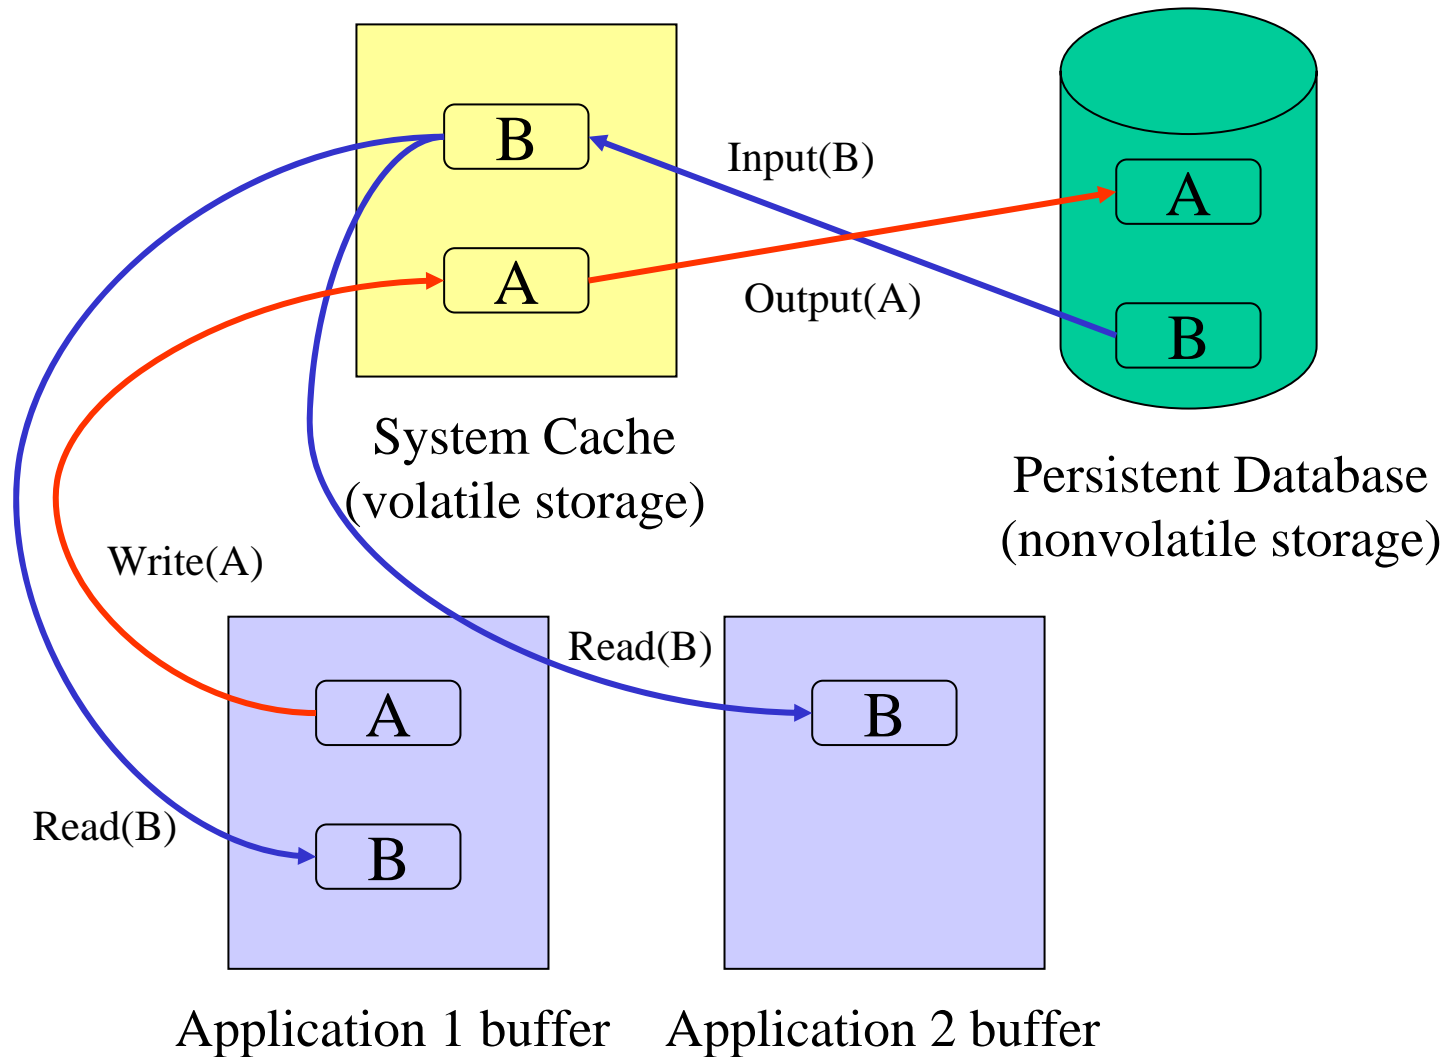
Storage Types

- Management of failure takes advantage of several types of storage.
- *Volatile storage*: does not survive system crashes
 - Main memory
 - Cache memory
- *Non-volatile storage*: survives system crashes
 - Disk
 - Tape
- *Stable storage*: a mythical form of storage that survives all failures
 - Approximated by maintaining multiple copies of distinct non-volatile media with independent failure modes.

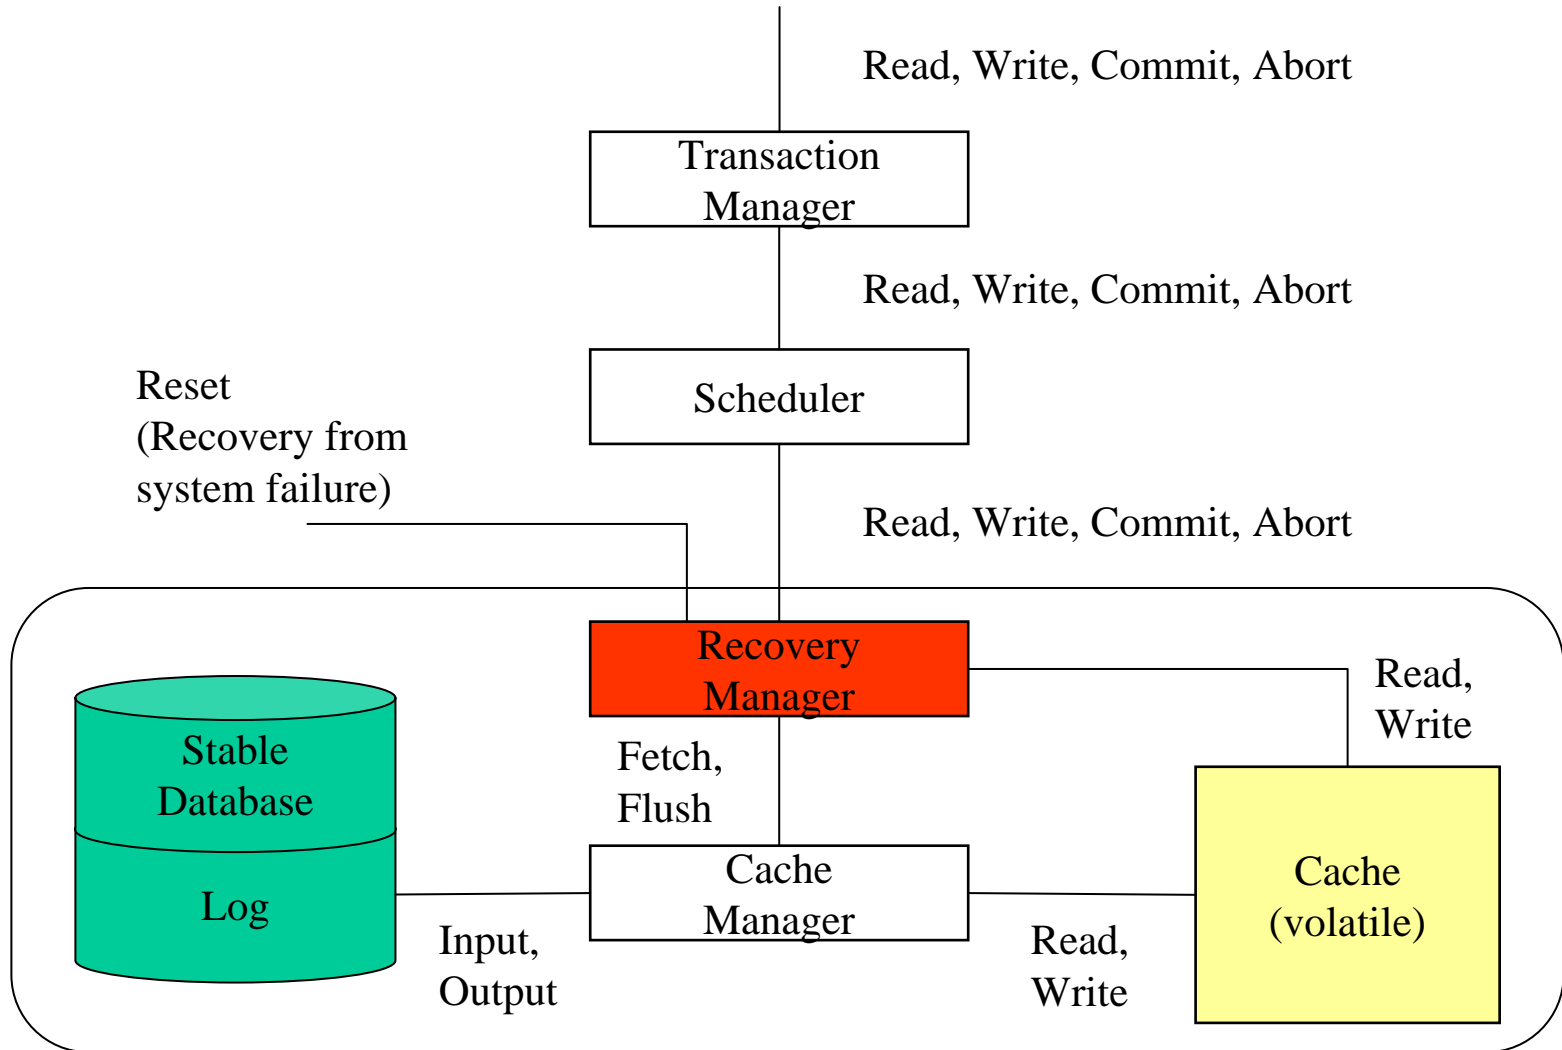
Storage Operations

- Transactions access and update the database.
- There are two operations for moving blocks with data items between disk and main memory.
 - Input(Q) transfers the block containing data item Q to main memory.
 - Output(Q) transfers the block containing Q to disk and replaces the appropriate block there.
- There are two operations for moving values between data items and application variables.
 - Read(Q,q) assigns the value of data item Q to variable q.
 - Write(Q,q) assigns the value of variable q to data item Q.
 - ◆ An Input(Q) may be necessary for both operations.

Movement of Values



System Model



Log-based Recovery

- During normal operation, the following occurs.
- Each transaction T when it starts, registers itself on the log, by outputting $\langle T \text{ starts} \rangle$.
- When a transaction T modifies a data item X by $\text{Write}(X)$, then the Recovery Manager first outputs an entry with the following fields to the log.
 - transaction's ID (i.e., T)
 - data item name (i.e., X)
 - old value of the item
 - new value of the item
- Then the Recovery Manager writes the new value of X into the cache.

Log-based Recovery, cont.

- Later, the Cache Manager can output the value to the portion of the database on disk.
- This allows output to stable storage to be asynchronous with application writes, which is especially effective in the presence of large main memories.
- When T reaches its last statement, the Recover Manager appends <T commits> to the log, and T then commits.
 - The transaction commits precisely when the commit entry (after all previous entries for this transaction) is outputted to the log.

Buffering of Log Entries and Data Items

- We have so far assumed that log entries are put on stable storage when they are created.
- Lifting this restriction, but imposing the following, leads to better performance.
 - Transaction T cannot commit before $\langle T \text{ commits} \rangle$ has been output to stable storage.
 - $\langle T \text{ commits} \rangle$ cannot be placed on stable storage before all other log entries pertaining to T are on stable storage.
 - A block of data items cannot be output to stable storage until all log entries pertaining to the data items are output.

Recovery From Failures

- When failures occur, the following operations that use the log are executed.
 - *Undo*: restore database to state prior to execution.
 - *Redo*: perform the changes to the database over again.
- Examples
 - Undo
 - ◆ Aborted transaction
 - ◆ Active transactions at time of crash
 - Redo
 - ◆ Committed transaction whose changes have not been placed in the database

Example

- We consider two transactions executed sequentially by the system.

■ T1:	Read(A)	T2:	Read(A)
	$A \leftarrow A + 50$		$A \leftarrow A + 10$
	Read(B)		Write(A)
	$B \leftarrow B + 100$		Read(D)
	Write(B)		$D \leftarrow D - 10$
	Read(C)		Read(E)
	$C \leftarrow 2C$		Read(B)
	Write(C)		$E \leftarrow E + B$
	$A \leftarrow A + B + C$		Write(E)
	Write (A)		$D \leftarrow D + E$
			Write(D)

- The initial values are:

$A=100, B=300, C=5, D=60, E=80$

Example, cont.

- The Log

1. <T1 starts>
2. <T1, B, old: 300, new: 400>
3. <T1, C, old: 5, new: 10>
4. <T1, A, old: 100, new: 560>
5. <T1 commits>
6. <T2 starts>
7. <T2, A, old: 560, new: 570>
8. <T2, E, old: 80, new: 480>
9. <T2, D, old: 60, new: 540>
10. <T2 commits>

A=100 B=300 C=5 D=60 E=80

T1:

Read(A)

$A \leftarrow A + 50$

Read(B)

$B \leftarrow B + 100$

Write(B)

Read(C)

$C \leftarrow 2C$

Write(C)

$A \leftarrow A + B + C$

Write (A)

T2:

Read(A)

$A \leftarrow A + 10$

Write(A)

Read(D)

$D \leftarrow D - 10$

Read(E)

Read(B)

$E \leftarrow E + B$

Write(E)

$D \leftarrow D + E$

Write(D)

- Output of B can occur anytime after entry 2 is output to the log, etc.

Example, cont.

- Assume a system crash occurs. The log is examined. Various actions are taken depending on the last instruction (actually) written on it.

1. <T1 starts>
2. <T1, B, ...>
3. <T1, C, ...>
4. <T1, A, ...>
5. <T1 commits>
6. <T2 starts>
7. <T2, A, ...>
8. <T2, E, ...>
9. <T2, D, ...>
10. <T2 commits>

Last Instruction	Action	Consequence
$I = 0$	Nothing	Neither T1 nor T2 has run
$1 \leq I \leq 4$	<i>Undo T1</i> : Restore the values of variables listed in 1-I old values	T1 has not run
$5 \leq I \leq 9$	<i>Redo T1</i> : Set the values of the variables listed in I-4 to values created by T1 <i>Undo T2</i> : Restore the values of variables listed in I-9 to those before T2 started execution	T1 ran T2 has not run
$I=10$	<i>Redo T1</i> <i>Redo T2</i>	T1 and T2 both ran

Recovery: Outline

- Logging Mechanism
- Recovery Algorithms
 - Undo/Redo
 - No-Undo/Redo
 - Undo/No-Redo
 - No-Undo/No-Redo
- Checkpoint
- Other Issues

The Undo/Redo Algorithm

- Following a failure, the following is done.
 - Redo all transactions for which the log has both “start” and “commit” entries.
 - Undo all transactions for which the log has “start” entry but no “commit” entry.
- Remarks:
 - In a multitasking system, more than one transaction may need to be undone.
 - If a system crashes during the recovery stage, the new recovery must still give correct results.
 - In this algorithm, a large number of transactions need to be redone, since we do not know what data items are on disk.

Undo/Redo, cont.

- Goal: Maximize efficiency during normal operation.
 - Some extra work is required during recovery time.
- Allows maximum flexibility of buffer manager.
 - Database outputs are entirely asynchronous, other than having to happen after the corresponding log entry outputs.
- Most complex at recovery time
 - Must implement both undo and redo.
- Also termed *immediate database modification*
- Note: requires before and after images in log or only after images if an initial before image is written to the log before first write.

No-Undo/Redo

- Algorithm
 - Don't output values to the disk until the commit log entry is on stable storage.
 - All writes go to the log and to the database cache.
 - Sometime after commit, the cached values are output to the disk.
- Advantages
 - Faster during recovery: no undo.
 - No before images needed in log.
- Disadvantages
 - Database outputs must wait.
 - Lots of extra work at commit time.
- Also called *deferred database modification*.

Undo/No-Redo


- Algorithm
 - All changed data items need to be output to the disk before commit.
 - ◆ Requires that the write entry first be output to the (stable) log.
 - At commit:
 - ◆ Output (*flush*) all changed data items in the cache.
 - ◆ Add commit entry to log.
- Advantages
 - No after images are needed in log.
 - No transactions need to be redone.
- Disadvantages
 - Hot spot data requires a flush for each committed write.
 - ◆ Implies lots of I/O traffic.

No-Undo/No-Redo


- Algorithm
 - No-undo → don't change the database during a transaction
 - No-redo → on commit, write changes to the database in a single atomic action
- Advantages
 - Recovery is instantaneous.
 - No recovery code need be written.
- Atomic databases writes of many pages is accomplished using the *shadow paging* technique.

No-Undo/No-Redo, cont.


- During a transaction:



- After preparing new directory for commit:



- After committing:



No-Undo/No-Redo Example, cont.

- Commit requires writing on disk of one bit.
- Restart is very fast: one disk read.
- Problems:
 - Access to stable storage is indirect (though directories may be possibly stored in main memory).
 - Garbage collection of stable storage is required.
 - Original layout of data is destroyed.
 - Concurrent transactions are difficult to support.


Recovery

- Logging Mechanism
- Recovery Algorithms
- Checkpoint
- Other Issues

Checkpoint

- Checkpoint speeds up recovery by flushing dirty pages to disk.
- During the execution in addition to the activities of the previous method, periodically perform checkpointing:
 1. Output the log buffers to the log.
 2. Force database buffers to the disk.
 3. Output an entry <checkpoint> on the log.
- During recovery
 - Undo all transactions that have not committed.
 - Redo all transactions that have committed *after* checkpoint.

Example - Recovery with Checkpoints



- $T1$ is ok.
- $T2$ and $T3$ are redone.
- $T4$ is undone.

Recovery

- Logging Mechanism
- Recovery Algorithms
- Checkpoint
- Other Issues
 - Media Failure and Archiving
 - Recovery Tuning

Media Failure

- Stable database or log may be wiped out.
 - Head crash
 - Operator error, e.g., mounting incorrect tape
 - Power lossage
 - Sabotage
 - Disgruntled employee
- Only solution: keep multiple copies on multiple devices with multiple failure rates.
- Mirroring
 - Each write applied to one disk, then to other.
 - On failure, use good device until bad device is fixed and current data restored.
 - ◆ During this time, there is no redundancy.
 - Increases the disk input rate, but not the disk output rate.

Archiving

- Protects against loss of non-volatile storage.
- Periodically dump entire database, in a consistent state, to an archive database.
- Mirror the logs.
- Recovery algorithms
 - Log destroyed: use other one; rebuild destroyed copy.
 - Stable database destroyed: apply log to archive.
 - Both logs destroyed: start again with archive.
 - ◆ Repeat lost transactions with information from paper records.
- Creating archive database can be expensive.
 - Many businesses are 24x7 operations.
 - Creating a consistent state is basically a huge read-only transaction.

Recovery Tuning

- Place the log on a separate disk.
 - This yields sequential log-writes and also improves reliability.
- Delay writes to the database disk(s).
 - Buffered database writes may reduce the random writes to the database disks.
 - Sometimes writes can be avoided completely (repeated updates of the same data item).
- Tune checkpoint and dump intervals.
 - Trade-off between performance of normal operation and performance of recovery
 - High-availability applications should do checkpoints very often, say every 20 minutes.

Interactive Transactions

- Interactive transactions are difficult to handle satisfactorily.
 - How to rollback a message to the user?
 - How to recall \$100 an automatic teller handed out?
- Correct but unsatisfactory solutions:
 - Forbid interactive transactions.
 - Send all messages after commit.
 - Break interactive transactions into smaller units of consistency that are transactions on their own.
- Another solution:
 - Define *compensating transactions*; this is not always possible.

Summary

- Undo/Redo
 - ◆ Before and after values written to log.
 - ◆ Writes occur whenever is convenient.
- No-Undo/Redo
 - ◆ Disk is written after commit.
 - ◆ No before values appear in log.
- Undo/No-Redo
 - ◆ Disk is written before commit.
 - ◆ No after values appear in log.
- No-Undo/Redo
 - ◆ Uses shadow paging.
 - ◆ Disk writes like undo/redo, except flush before commit record.
- Checkpoint
 - ◆ Flush buffers before writing checkpoint record