

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic SQL Queries
- More Powerful SQL Queries
- Constraints in SQL

History of SQL

- *SEQUEL*: Structured English Query Language
 - 23-page research report in 1974 from IBM San Jose Research Labs
 - System R prototype
- *SEQUEL 2*: A Unified Approach to Data Definition, Manipulation and Control, IBM Systems Journal, 1976
- *SQL*: Structured Query Language (IBM's DB2 in 1983)

Standardization of SQL

- SQL-86
 - ISO 9075-1987, “Database Language SQL”
 - A somewhat cleaned up version of DB2 SQL
- SQL-89
 - Added embedded SQL
 - 150 pages
- SQL-92 (also called SQL2)
 - ISO/IEC 9075:1992
 - Many new DDL and DML features
 - 500 pages
- SQL:1999/SQL:2003 (also called SQL3)
 - Object-oriented features
 - Fifteen parts
 - ~5000 pages

Levels of SQL-92

- Entry SQL
 - SQL-89 with some small changes
- Intermediate SQL
 - Approximately half of the new features of SQL-92
 - Expected that initial implementations will be at this level.
- Full SQL
 - All features in the standard
 - Few DBMS supports this level even today
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
 - Not all examples here may work on your particular system.

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic SQL Queries
- More Powerful SQL Queries
- Constraints in SQL

SQL's Data Model

- Tables are multi-sets (bags) of rows.
- Each row of a table has the same columns, over the same domains.
- Duplicates are allowed
 - unless explicitly disallowed via **UNIQUE**.
- Order in which tuples are stored cannot be specified
 - But some DBMS vendors allow this in various ways

Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *n* digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- More are covered in Chapter 4.

Data Definition in SQL

- Three statements are used to define the schema in SQL.
 - **CREATE**
 - **DROP**
 - **ALTER**
- These statements apply to
 - Tables (and indirectly columns)
 - Views
 - Domains
 - Assertions
 - Character sets
 - ◆ Internationalization of software, hard-work and boring
 - Translations (between character sets)

CREATE TABLE

- Specifies a new base table.
- Name
- Columns with
 - Name
 - Data type
 - Column constraints
 - Default value
- Table constraints
 - Primary and unique keys
 - Foreign keys

Video Store Example

```
CREATE TABLE Film (  
    FilmID          NUMERIC (5),  
    Title           VARCHAR (50),  
    PubDate         DATE ,  
    RentalPrice     NUMERIC (5,2),  
    Distributor     VARCHAR (20),  
    Kind            CHAR (1),  
    RecommededAge   NUMERIC (2),  
    SpokenLanguage  VARCHAR (15),  
    SubtitleLanguage VARCHAR (15))
```

```
CREATE TABLE Reserves (  
    CustomerID      NUMERIC (5),  
    FilmID          NUMERIC (5),  
    ResDate         DATE )
```

```
CREATE TABLE Customer (  
    CustomerID      NUMERIC (5),  
    Name            VARCHAR (50),  
    Street          VARCHAR (50),  
    City            VARCHAR (50),  
    State           CHAR (2))
```

Column Defaults

- If not specified, the value of the column will be **NULL**.
- A particular value can also be specified.

```
CREATE TABLE Film (  
    ...  
    SpokenLanguage VARCHAR (15) DEFAULT 'English',  
    ...  
)
```

```
CREATE TABLE Reserves (  
    ...  
    ResDate DATE DEFAULT CURRENT_DATE,  
    ...  
)
```

Column Constraints

- **NOT NULL**

```
CREATE TABLE Film ( ...  
                    PubDate DATE NOT NULL, ... )
```

- If no default is given, then **INSERT** must specify a value, otherwise the constraint will be violated.

- **UNIQUE**

```
CREATE TABLE Customer ( ...,  
                        ID NUMERIC (5) UNIQUE, ... )
```

- Ensures that there are no duplicate tuples, with the same column value.
- Can have one value that is **NULL**, unless **NOT NULL** is specified.

Column Constraints, cont.

- **PRIMARY KEY**

```
CREATE TABLE Film (  
    FilmID NUMERIC (5) PRIMARY KEY,  
    ...)
```

- Implies **NOT NULL, UNIQUE** (entity integrity).

- **CHECK** (*predicate*)

```
CREATE TABLE Film ( ...,  
    Kind CHAR(1) CHECK (Kind IN ('F', 'M', 'E')),  
    ...)
```

- Any predicate that may occur in an SQL where clause is allowed.
- The predicate is assumed to range only over the row in which the constraint appears.

DROP TABLE

- Used to remove a table and its definition
- Once dropped, the table can no longer be used in queries, updates, or any other commands since its description no longer exists.
- Example: remove the file table created before

```
DROP TABLE Film
```

ALTER TABLE

- Used to
 - Add a column
 - Drop a column
 - Change a column's default
 - Add a constraint
 - Drop a constraint

- Example

```
ALTER TABLE Film ADD COLUMN PurchasePrice  
                                NUMERIC (5,2)
```

- The new attribute will have **NULLs** in all the tuples of the table immediately after the command is executed.
- Hence, the **NOT NULL** constraint is not allowed for such an attribute, unless a **DEFAULT** is specified.
- The database users must still enter a value for the new column `PurchasePrice` for each `Film` row.

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic SQL Queries
- More Powerful SQL Queries
- Constraints in SQL

Data Manipulation in SQL

- Basic SQL statements for data manipulation
 - **SELECT** - retrieve data from the a database
 - **INSERT** - insert data into a table
 - **UPDATE** - updates existing data within a table
 - **DELETE** - deletes all records from a table, but the space for the records remain
 - ◆ Different from DROP table

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic SQL Queries
 - **SELECT-FROM-WHERE** Statement
 - **NULL** Values
 - **ORDER BY**
 - **GROUP BY** and **HAVING**
- More Powerful SQL Queries
- Constraints in SQL

The Basic SQL Statement

- The clauses are specified in the following order.
 - **SELECT** *column(s)*
 - **FROM** *table list*
 - **WHERE** *condition*
 - **GROUP BY** *grouping column(s)*
 - **HAVING** *group condition*
 - **ORDER BY** *sort list*

Retrieval Queries in SQL

- SQL has one basic statement for retrieving information from a database; the **SELECT** statement.
- This is not the same as the select (σ) operation of the relational algebra.
- The basic form of the SQL **SELECT** statement is called a *mapping* or a *select-from-where block*.

```
SELECT column list  
FROM    table list  
WHERE   condition
```

- Equivalent to the following relational algebra formula.

$$\pi_{column\ list}(\sigma_{condition}(table_1 \times \dots \times table_n))$$

Queries Over Several Tables

- List the IDs of the films that are expensive and have been reserved.

```
SELECT Film.FilmID
FROM    Film, Reserves
WHERE    RentalPrice > 4
           AND Film.FilmID = Reserves.FilmID
```

- List the customer names who have reserved a film.

```
SELECT DISTINCT Name
FROM            Customer, Reserves
WHERE           Reserves.CustomerID = Customer.CustomerID
```

Queries Over Several Tables, cont.

- List the customer names who have reserved an expensive film.

```
SELECT DISTINCT Name
FROM    Customer, Film, Reserves
WHERE   Reserves.CustomerID = Customer.CustomerID
          AND Reserves.FilmID = Film.FilmID
          AND RentalPrice > 4
```

- List the streets of customers who have reserved foreign films.

```
SELECT Street
FROM    Customer, Film, Reserves
WHERE   Reserves.CustomerID = Customer.CustomerID
          AND Reserves.FilmID = Film.FilmID
          AND Kind = 'F'
```

NULL Values

- **NULL** values are used if
 - a value is missing
 - a value is unknown
 - a value does not exist
 - ...(29 possible interpretations for **NULL** were identified)
- Operations and functions return **NULL** if one of their values in **NULL**.
 - Aggregates is an exception
 - ◆ they ignore **NULL** values

NULL Values, Cont.

- In predicates, **NULL** is handled with a three-valued logic approach, i.e., a logic with truth values TRUE, FALSE, and UNKNOWN.

Value of x	Condition	Result
NULL	X IS NULL	TRUE
NULL	X IS NOT NULL	FALSE
10	X IS NULL	FALSE
10	X IS NOT NULL	TRUE
NULL	X = NULL	UNKNOWN
NULL	X <> NULL	UNKNOWN
10	X = NULL	UNKNOWN
10	X <> NULL	UNKNOWN

ORDER BY

- Can sort the result of a select, using **ORDER BY**.

- Who has reserved a film?

```
SELECT Name
FROM   Customer, Reserves
WHERE  Customer.CustomerID = Reserves.CustomerID
ORDER BY Name
```

- Can also sort in descending order, via **DESC**

- (**ASC** is the default).

```
SELECT Name
FROM   Customer, Reserves
WHERE  Customer.CustomerID = Reserves.CustomerID
ORDER BY Name DESC
```

- Only columns in the select list can be used for the ordering.

SELECT in the FROM Clause

- The table in a from clause can itself be a select statement.
- List the customers who have reserved an expensive film.

```
SELECT CustomerID
FROM    Reserves, (SELECT *
                   FROM    Film
                   WHERE    RentalPrice > 4) AS Expensive
WHERE    Reserves.FilmID = Expensive.FilmID
```

- A correlation name is required in such case
 - Expensive in this case
 - Use **AS** as the keyword of naming
 - ◆ **AS** should be removed on Oracle

Aggregates

- Aggregates operate on the set of values of a column of a table, and return a single value.
- SQL offers 5 built-in aggregate functions:
 - **SUM**: sum of values
 - **AVG**: average value
 - **MAX**: maximum value
 - **MIN**: minimum value
 - **COUNT**: number of values
- What is the total rental price of all films?

```
SELECT SUM (RentalPrice)  
FROM Film
```

Aggregates, cont.

- What is the average rental price of reserved films?

```
SELECT AVG(RentalPrice)
FROM   Reserves, Film
WHERE  Reserves.FilmID = Film.FilmID
```

- How many reservations are there?

```
SELECT COUNT ( * )
FROM   Reserves
```

- **DISTINCT** can eliminate duplicates before computing the aggregate.

- How many cities do customers reside in?

```
SELECT COUNT (DISTINCT City)
FROM   Customer
```

- **DISTINCT** cannot be used with `COUNT (*)`

GROUP BY

- The aggregate can be applied to several groups by specifying **GROUP BY**.
- What is the average rental price of all films, by kind?

```
SELECT DISTINCT Kind, AVG(RentalPrice)
FROM Film
GROUP BY Kind
```

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3
Rocky V	Action	2



<i>Kind</i>	<i>AvgRentalPrice</i>
Action	2.5
Western	3.5
Fiction	3

A Special GROUP BY Example

Films:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3
Rocky V	Action	2

SELECT *
FROM Films
GROUP BY Kind

Result:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Rocky V	Action	2
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3

HAVING

- Individual groups can be eliminated by using **HAVING** that follows **GROUP BY**.
- List the customers whose average rental price for reserved films is greater than \$3.

```
SELECT Name
FROM    Customer, Reserves, Film
WHERE    Customer.CustomerID = Reserves.CustomerID
           AND Reserves.FilmID = Film.FilmID
GROUP BY Name
HAVING AVG(RentalPrice) > 3
```

HAVING Example

Films:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3
Rocky V	Action	2

FROM Films
GROUP BY Kind
HAVING Kind = 'Western'
OR Kind = 'Action'

Result:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Rocky V	Action	2
Unforgiven	Western	4
Once upon a Time in the West	Western	3

HAVING Example (cont)

Films:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3
Rocky V	Action	2



```
FROM Films  
GROUP BY Kind  
HAVING MAX(RentalPrice) <= 3
```

Result:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Rocky V	Action	2
Star Wars	Fiction	3

HAVING Example (cont)

Films:

<i>Title</i>	<i>Kind</i>	<i>RentalPrice</i>
Lethal Weapon 4	Action	3
Unforgiven	Western	4
Once upon a Time in the West	Western	3
Star Wars	Fiction	3
Rocky V	Action	2

FROM Films
GROUP BY Kind
HAVING RentalPrice <= 3

Must have an aggregate function!

ILLEGAL!

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic SQL Queries
- More Powerful SQL Queries
 - Joins
 - Modifications
 - Subqueries
 - IN/ALL/ANY/EXISTS
 - Views
- Constraints in SQL

Joins

- Inner Joins
 - Cartesian product
 - Old style join
 - ◆ Specified by including several tables in the from clause.
 - ◆ Can also be specified directly: **INNER JOIN**.
 - Condition join
 - ◆ Specified by including a predicate in the where clause.
 - Cross join
 - Natural join
 - Column name join
- Outer Join
- Union Join

Cross Join

- The *cross join* is equivalent to a Cartesian product.
- List all the customer information, including all reservations.

```
SELECT *  
FROM    Customer CROSS JOIN Reserves
```

- Note that the result will include too much information, since all reservations will be paired with all customers.
- This query is equivalent to the following.

```
SELECT Customer.*, Reserves.*  
FROM    Customer, Reserves
```

Natural Join

- As with the relational algebra, the natural join in SQL eliminates duplicate columns.
- List all the customer information, and all their reservations.

```
SELECT *  
FROM    Customer NATURAL JOIN Reserves
```

- This query retrieves the correct information.
- It is equivalent to the following.

```
SELECT Customer.*, FilmID, ResDate  
FROM    Customer, Reserves  
WHERE    Customer.CustomerID = Reserves.CustomerID
```

Column Name Join

- The *column name join* uses only some of the columns that a natural join would use in the equality test.
- The `Customer` and `Reserves` tables both include address information. Say we wish to join only on the `CustomerID` column, which is also present in both.

```
SELECT *  
FROM    Customer JOIN Reserves USING (CustomerID)
```

Outer Join

- In the **FROM** list
 - **LEFT OUTER JOIN ON** *predicate*
 - **RIGHT OUTER JOIN ON** *predicate*
 - **FULL OUTER JOIN ON** *predicate*
- List the films, along with the customers who reserved each, if applicable.

```
SELECT Title, Name
FROM    Film LEFT OUTER JOIN Reserves
          ON Film.FilmID = Reserves.FilmID, Customer
WHERE    Customer.CustomerID = Reserves.CustomerID
```


Modifications

- There are three modification statements.
 - **INSERT**
 - **UPDATE**
 - **DELETE**
- For insertions, either values can be specified, or a select statement provides the values.
- Enter a reservation for Eric for the film 332244.

```
INSERT INTO Reserves
```

```
    VALUES (123456, 332244, CURRENT_DATE)
```

- Let Melanie reserve all the films that Eric has reserved.

```
INSERT INTO Reserves
```

```
    SELECT C2.CustomerID, FilmID, CURRENT_DATE
```

```
    FROM    Reserves, Customer AS C1, Customer AS C2
```

```
    WHERE    C1.Name = 'Eric'
```

```
            AND C1.CustomerID = Reserves.CustomerID
```

```
            AND C2.Name = 'Melanie'
```

DELETE

- A where clause identifies which rows to remove from the table.
- Delete all the reservations of customer 123456.

```
DELETE FROM Reserves  
WHERE CustomerID = 123456
```

- Other tables can be consulted to determine which rows should be removed.
- Delete all of Eric's reservations.

```
DELETE FROM Reserves  
WHERE CustomerID = ANY(SELECT CustomerID  
                        FROM Customer  
                        WHERE Name = 'Eric')
```

Someone from the set 

A set is returned 

UPDATE

- Increase the rental price of all films by 10%.

```
UPDATE Film
      SET RentalPrice = RentalPrice * 1.10
```

- The update statement has an optional where clause.
- Increase the rental price of foreign films by 10%.

```
UPDATE Film
      SET RentalPrice = RentalPrice * 1.10
      WHERE Kind = 'F'
```

SELECT: Subqueries

- A SELECT may be nested

SELECT ...

FROM ...

**WHERE <cond> (SELECT ...
FROM ...
WHERE ...)**

- Subqueries may produce
 - A scalar (single value)
 - A single--column table
 - ◆ **ANY, ALL, IN, EXISTS**
 - A multiple-column table
 - ◆ **EXISTS**
- Correlated subqueries

Scalar Producing Subquery

- The subquery produces a single value that can be compared
- What are the IDs of customers with the same name as the customer with ID 123456?

```
SELECT CustomerID
FROM Customer
WHERE name = (SELECT name
                FROM Customer
                WHERE CustomerID = 123456)
```

Single Attribute Producing Subquery

- The subquery produces a table with a single column
- **IN**
 - true if value exists in result of subquery
- *comparisonOperator* **ANY**
 - true for comparison with at least one tuple in subquery produced table
- *comparisonOperator* **ALL**
 - true for comparison with every tuple in subquery produced table

IN

- **IN** is equivalent to a restricted form of exists:

$$V \text{ IN } r \Leftrightarrow \exists t \in r (t = V)$$

- (246800 **IN**

123456
246800
369121

) is true.

- (333333 **IN**

123456
246800
369121

) is false.

- (333333 **NOT IN**

123456
246800
369121

) is true.

IN Query Example

- List the IDs of the films that are expensive and have been reserved.

```
SELECT FilmID
FROM    Film
WHERE    RentalPrice > 4
          AND FilmID IN (SELECT FilmID
                          FROM    Reserves)
```

```
SELECT FilmID
FROM    Reserves
WHERE    FilmID IN (SELECT FilmID
                    FROM    Film
                    WHERE RentalPrice > 4)
```


IN, cont.

- List the IDs of the expensive films that have not been reserved.

```
SELECT FilmID
FROM    Film
WHERE    RentalPrice > 4
           AND FilmID NOT IN ( SELECT FilmID
                                FROM    Reserves )
```

ANY

- **ANY** is also equivalent to exists:

$$V \text{ comp } \mathbf{ANY} \ r \Leftrightarrow \exists t \in r \ (V \text{ comp } t)$$

- $(246800 < \mathbf{ANY} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is true.

- $(369121 < \mathbf{ANY} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is false.

ANY, cont.

- $(246800 = \mathbf{ANY} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is true.

- $(246800 <> \mathbf{ANY} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is true.

- Comparison with **IN**

- $(= \mathbf{ANY}) \iff \mathbf{IN}$
- $(<> \mathbf{ANY}) \iff \mathbf{NOT IN}$

ANY Query Example

- Which films rent for more than *some* foreign film?

```
SELECT Title
FROM   Film
WHERE  RentalPrice > ANY( SELECT RentalPrice
                           FROM   Film
                           WHERE  Kind = 'F' )
```

ALL

- **ALL** is equivalent to for all:

$$V \text{ comp } \mathbf{ALL} \ r \Leftrightarrow \forall t \in r (V \text{ comp } t)$$

- $(246800 < \mathbf{ALL} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is false.

- $(100000 < \mathbf{ALL} \begin{array}{|c|} \hline 123456 \\ \hline 246800 \\ \hline 369121 \\ \hline \end{array})$ is true.

ALL, cont.

- (246800 = **ALL**

123456
246800
369121

) is false.

- (246800 <> **ALL**

123456
246800
369121

) is false.

- Comparison with **IN**

- (<> **ALL**) \Leftrightarrow **NOT IN**
- (= **ALL**) \nLeftrightarrow **IN**

ALL Query Example

- Which films rent for more than *all* foreign films?

```
SELECT Title
FROM Film
WHERE RentalPrice > ALL( SELECT RentalPrice
                          FROM Film
                          WHERE Kind = 'F' )
```

ALL Query Exmample, cont.

- Find the film(s) with the highest rental price.

```
SELECT Title
FROM   Film
WHERE  RentalPrice >= ALL (SELECT RentalPrice
                           FROM   Film )
```


Correlated Subqueries

- Subquery must be re-evaluated for each tuple in outer **SELECT**
- Table variable is used
- Find the customers who live at more than one address
 - assume just street name needs to be different.

```
SELECT Name
FROM    Customer C
WHERE    CustomerID IN
           ( SELECT D.CustomerID
             FROM    Customer D
             WHERE    C.Street <> D.Street );
```

Binding

- SQL follows binding rules from tuple relational calculus.
- Previous query can be expressed without correlation name D (using a default correlation name of Customer).

```
SELECT Name
FROM    Customer C
WHERE    CustomerID IN
           ( SELECT CustomerID
             FROM Customer
             WHERE C.Street <> Street );
```

Multiple Attribute Producing Subquery

- The subquery produces a table with several columns
- **EXISTS**
 - true if subquery produced table has a tuple
- **NOT EXISTS**
 - true if subquery produced table is empty

EXISTS

$$\text{EXISTS } T \Leftrightarrow T \neq \emptyset$$

- List the customers who live in Pullman and have a film reserved.

```
SELECT Name
FROM Customer
WHERE City = 'Pullman'
      AND EXISTS( SELECT *
                  FROM Reserves
                  WHERE Reserves.CustomerID =
                      Customer.CustomerID)
```

EXISTS, cont.

- Often, **EXISTS** can be replaced with another correlation name.
- List the customers who live in Pullman and have a film reserved.

```
SELECT Name
FROM    Customer
WHERE   City = 'Pullman'
          AND CustomerID IN (SELECT CustomerID
                              FROM    Reserves)
```

```
SELECT Name
FROM    Customer, Reserves
WHERE   City = 'Pullman'
          AND Customer.CustomerID = Reserves.CustomerID
```

NOT EXISTS

$$\text{NOT EXISTS } T \Leftrightarrow T = \emptyset$$

- List the customers who live in Pullman but have no films reserved.

```
SELECT Name
FROM Customer
WHERE City = 'Pullman'
      AND NOT EXISTS (SELECT *
                      FROM Reserves
                      WHERE Reserves.CustomerID =
                          Customer.CustomerID)
```

NOT EXISTS, cont.

- Often, **NOT EXISTS** can be replaced with **NOT IN**.
- List the customers who live in Pullman but have no films reserved.

```
SELECT Name
FROM    Customer
WHERE    City = 'Pullman'
AND      CustomerID NOT IN ( SELECT CustomerID
                                FROM    Reserves )
```

View Definition

- *Views* provide a mechanism to hide certain data from a certain class of users.
- To create a view we use the command:

CREATE VIEW *name* **AS** *query expression*

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- When a view is created, the query expression is stored in the database; the expression is substituted into queries using the view.

View Example

- Consider a person who needs to know only customer names and films they've reserved.
- Define a view of all customers holding reservations, and the films they have reserved.

```
CREATE VIEW Reservations AS  
  SELECT Name, Title  
  FROM    Customer, Reserves, Film  
  WHERE   Customer.CustomerID =  
           Reserves.CustomerID  
           AND Reserves.FilmID = Film.FilmID
```

Query on View

- A query on a view is transformed into a query on the base tables.

- Query on view `Reservations`:

What films has Melanie reserved?

```
SELECT Title
FROM    Reservations
WHERE    Name = 'Melanie'
```

- Query on the base tables

```
SELECT Title
FROM    Customer, Reserves, Film
WHERE    Customer.CustomerID = Reserves.CustomerID
          AND Reserves.FilmID = Film.FilmID
          AND Name = 'Melanie'
```

SQL: Outline

- Brief History
- Data Definition Language
- Data Manipulation Language
- Basic Structure of SQL Queries
- More Powerful SQL Queries
- Constraints in SQL
 - Column Constraints (We've seen them in CREATE TABLE)
 - Table Constraints
 - Referential Integrity
 - Domain Constraints
 - Assertion

Table Constraints

- **UNIQUE** (*column names*)

```
CREATE TABLE Reserves ( ...,  
    UNIQUE ( CustomerID, FilmID, ResDate )  
)
```

- **PRIMARY KEY**

```
CREATE TABLE Reserves ( ...,  
    PRIMARY KEY ( CustomerID, FilmID, ResDate )  
)
```

- Requires that all columns of the primary key be **NOT NULL**.

- **FOREIGN KEY** (*column(s)*)

REFERENCES *table* [(*column(s)*)]

```
CREATE TABLE Reserves ( ...,  
    FOREIGN KEY (CustomerID) REFERENCES Customer)
```

Table Constraints: **CONSTRAINT / CHECK**

- **CHECK** (*predicate*)


```
CREATE TABLE Film (... ,  
    CONSTRAINT (  
        NOT (SpokenLanguage = SubtitleLanguage)))
```

- The predicate may include nested select statements, mentioning the current or other tables.

```
CREATE TABLE CheckedOut (... ,  
    CHECK ( 'OK' IN  
        ( SELECT Status  
          FROM   VideoTape  
          WHERE  FilmID = VideoTape.FilmID)), ... )
```

Referential Integrity

- Referential integrity says “pointed to” information must exist.
 - A foreign key points to data in some relation
- Example
 - Customer information must exist for a customer to reserve a film.
 - No CustomerID can be in Reserves and not in Customer.
- Can be specified as a column constraint

```
CREATE TABLE Reserves (...  
    CustomerID INTEGER  
        CONSTRAINT ReservesToCustomerFK  Constraint name  
        REFERENCES Customer(CustomerID),  
    ... )
```

- Can be specified as a table constraint

```
CREATE TABLE Reserves (... ,  
    CONSTRAINT ReservesToCustomerFK  
        FOREIGN KEY (CustomerID) REFERENCES  
        Customer(CustomerID) ... )
```

Referential Integrity Violation Remedies

- Can specify **ON UPDATE** and **ON DELETE** options
- Example

```
CREATE TABLE Reserves (... ,  
    CONSTRAINT ReservesToCustomerFK  
    FOREIGN KEY (CustomerID) REFERENCES  
        Customer(CustomerID)  
    ON DELETE CASCADE ON UPDATE SET NULL  
    ... )
```

- Options (next slides)
 - Note: *Child* table - has the foreign key, references key in *parent* table

Example: Customer is parent, Reserves is child.

Remedy Options (1)

- None
 - Update or delete parent value
 - No change to matching child value
- Restrict
 - Cannot update or delete parent value if one or more matching values exist in the child table
 - No change to matching child value

Restrict Remedy Example

- Update Customer with ID 2, changing ID to 5

<

- Cannot update. Must first drop tuple from Reserves.

Remedy Options (2)

- Cascade
 - Update or delete parent value
 - Update or delete matching values in child table
- Set Null
 - Update or delete parent value
 - Set matching values in child table to NULL
- Set Default
 - Update or delete parent value
 - Set matching values in child table to default value

Cascade Remedy Example

- Delete Customer with ID 2

Before

Customer

Name	ID
Fred	2
Pam	4
Fred	3

Reserves

FilmID	CID
7	2
2	4

After

Customer

Name	ID
Pam	4
Fred	3

Reserves

FilmID	CID
2	4

Cascade Remedy Example

- Update Customer with ID 2, changing ID to 5

Before

Customer

Name	ID
Fred	2
Pam	4
Fred	3

Reserves

FilmID	CID
7	2
2	4

After

Customer

Name	ID
Fred	5
Pam	4
Fred	3

Reserves

FilmID	CID
7	5
2	4

Set Null Remedy Example

- Delete Customer with ID 2

Before

Customer

Name	ID
Fred	2
Pam	4
Fred	3

Reserves

FilmID	CID
7	2
2	4

After

Customer

Name	ID
Pam	4
Fred	3

Reserves

FilmID	CID
7	NULL
2	4

Domain Constraints

- No type constructors exist in SQL-92. SQL3 includes abstract data types.
- SQL-92 allows domains to be defined. These pull together a specific data type, as well as other characteristics of the type.
 - Size
 - Default
 - Constraints
- Example

```
CREATE DOMAIN CustomerDomain INT (6)
CHECK ( VALUE IS NOT NULL
        AND VALUE > 99999)
```


Assertions

- An assertion is a standalone constraint that we wish the database ALWAYS to satisfy. Such a restriction normally affects more than one table.
 - Once created, it is checked for validity immediately.
 - Any future modification on the database will only allowed if it does not violates the assertion. (Potential high overhead!)
- No one can reserve more than 3 films.

```
CREATE ASSERTION res_limit
CHECK ( NOT EXISTS (
    SELECT Name
FROM      Customer
WHERE      3 < ( SELECT COUNT ( * )
                  FROM      Reserves
                  WHERE      Reserves.CustomerID =
                           Customer.CustomerID ) ) )
```

Summary

- The SQL Query Language has two main parts
 - DDL statements
 - ◆ **CREATE, DROP, ALTER**
 - ◆ Constraint definition
 - DML statements
 - ◆ **SELECT .. FROM .. WHERE ..**
 - ◆ **INSERT, DELETE, UPDATE**
 - ◆ Impacts of constraints
- A query in SQL can consist of up to six clauses, but only the first two, **SELECT** and **FROM**, are mandatory.
- Complex SQL statement can be build using
 - **UNION/INTERSECT/EXCEPT** (not covered in these slides)
 - **IN/ALL/ANY/EXISTS**