

# Transactions: Outline

---

- The Transaction Concept
  - ACID Properties
  - Transaction States
- The Concurrency Problem
  - System Model
  - Schedules
- Serializability of Schedules
  - View Serializable
  - Conflict Serializable
- Recoverability
  - Recoverable Schedules
  - Cascadeless Schedules

# Transactions

---

- The Transaction Concept
  - ACID Properties
  - Transaction States
- The Concurrency Problem
- Serializability of Schedules
- Recoverability

# Transaction Concept

---

- A **transaction** is a *unit* of program execution that *accesses* and possibly *updates* various data items.
- A transaction must see a *consistent* database.
- During transaction execution the database may be *temporarily inconsistent*.
- When the transaction completes successfully (is committed), the database must be consistent.
- After a transaction commits, the changes it has made to the database persist, even if there are system failures.
- Multiple transactions can execute in parallel.
- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions

# ACID Properties

---

- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

# Example of Fund Transfer

---

- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- **Atomicity requirement** — if the transaction fails after step 3 and before step 6, the system should ensure that its updates are not reflected in the database, else an inconsistency will result.
- **Consistency requirement** – the sum of A and B is unchanged by the execution of the transaction.

# Example of Fund Transfer, cont.

---

- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- **Isolation requirement** — if between steps 3 and 6, another transaction is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be).
  - Isolation can be ensured trivially by running transactions **serially**, that is one after the other.
  - However, executing multiple transactions concurrently has significant benefits, as we will see later.
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist despite failures.

# Transaction Implementation

---

- A transaction is a *unit* of program execution.
- Transaction boundaries are user-defined.
  - **COMMIT** work
  - **ROLLBACK** work
    - ◆ Abort the transaction.
- Implementation of ACID properties
  - Atomicity: typically implemented via logs
  - Consistency: according to constraints/checks/assertions
  - Isolation: typically implemented via locks
  - Durability: typically implemented via logs

# Transaction States

---

- **Active** – the initial state; the transaction stays in this state while it is executing.
- **Partially committed** – after the final statement has been executed.
- **Failed** -- after the discovery that normal execution can no longer proceed.
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.

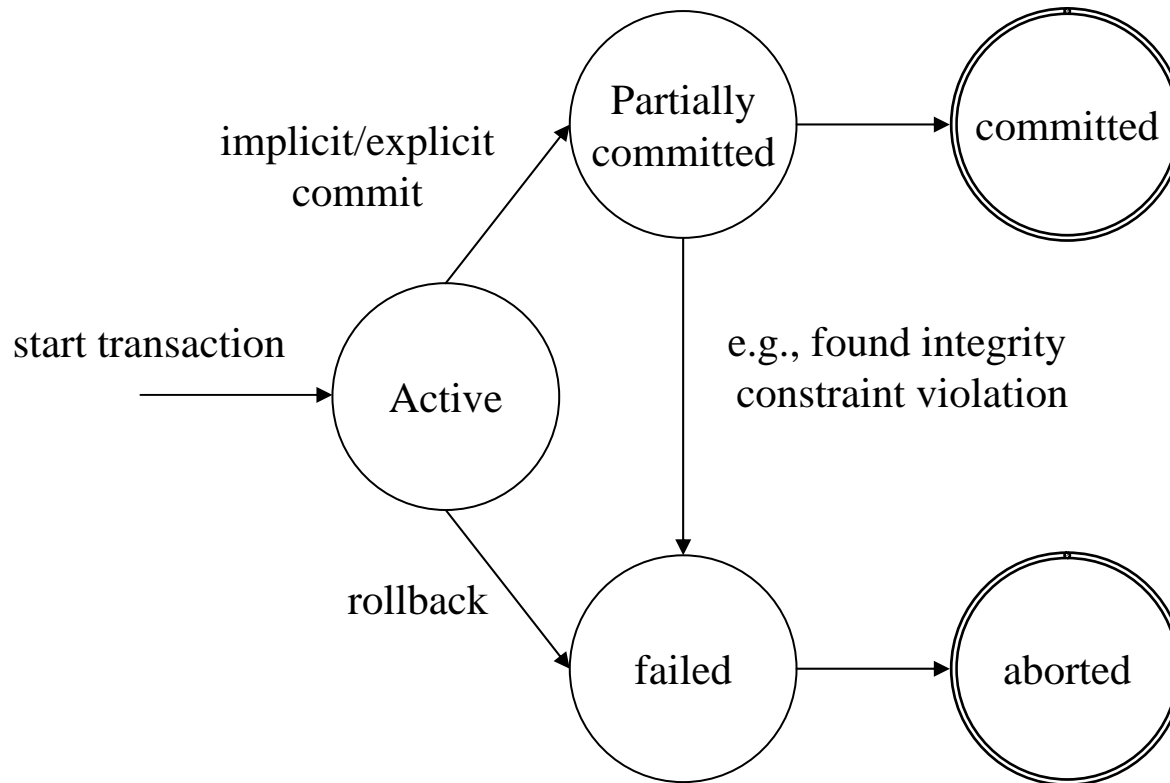
Two options after it has been aborted:

- restart the transaction; can be done only if no internal logical error
- kill the transaction
- **Committed** – after successful completion.



# Transaction State Diagram

---

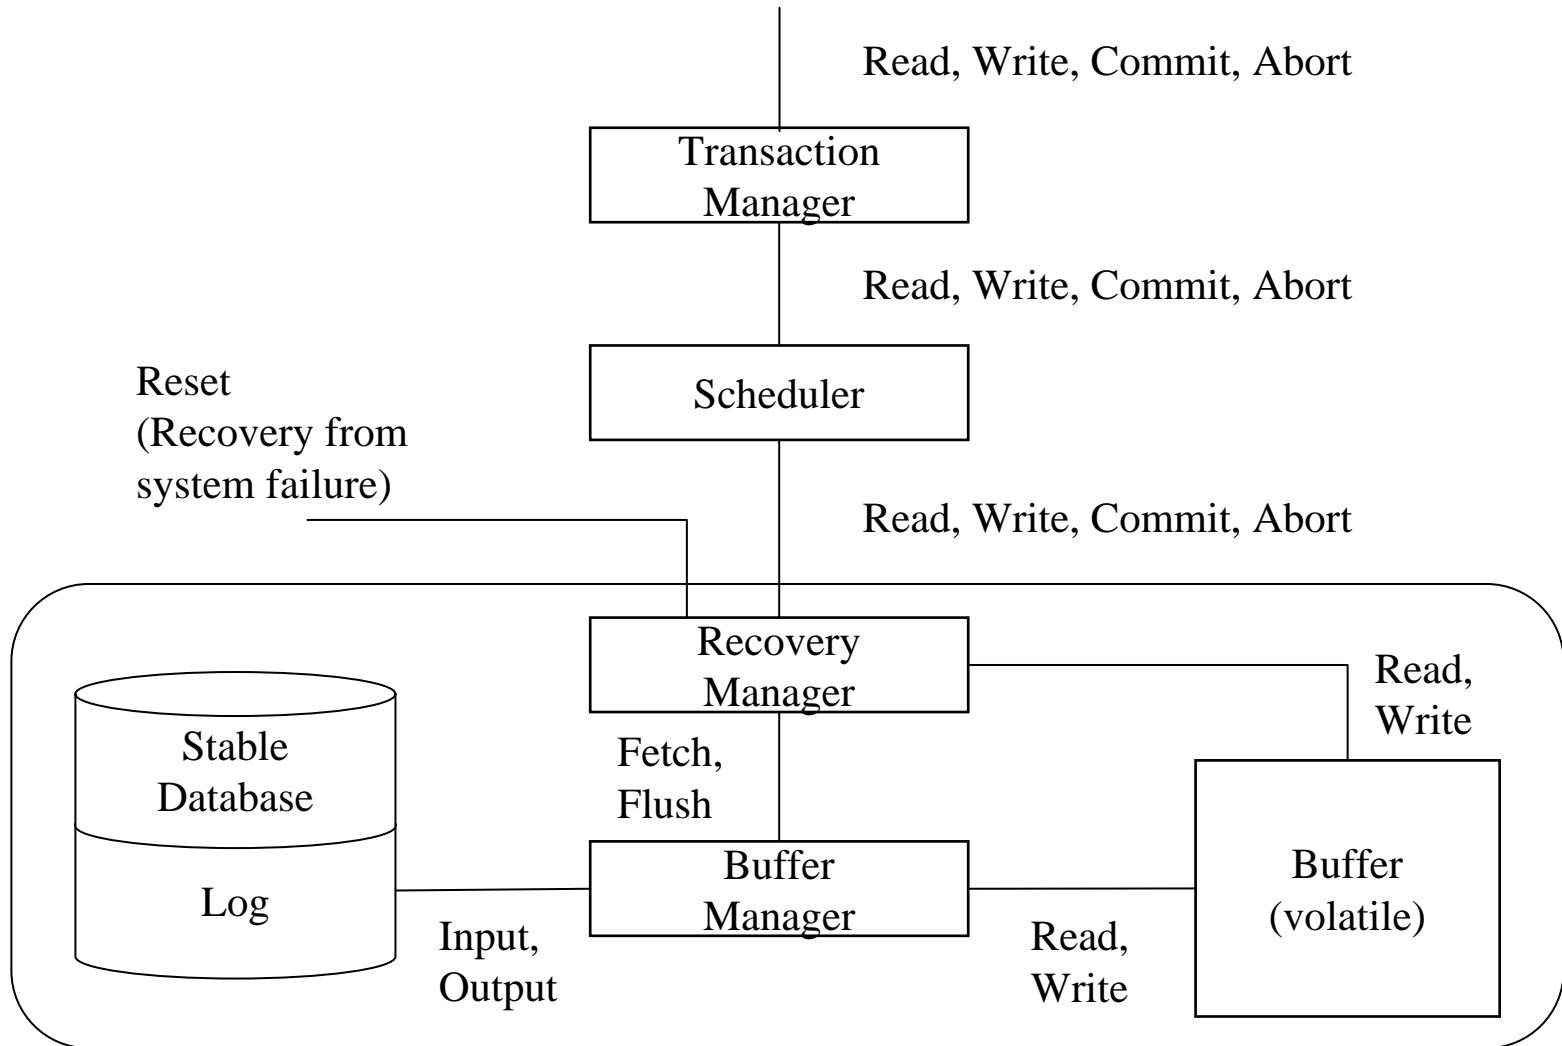


# Transactions

---

- The Transaction Concept
- The Concurrency Problem
  - System Model
  - Schedules
- Serializability of Schedules
- Recoverability

# System Model



# The Concurrency Problem

---

- Model: Centralized system with concurrent access by several users.
  - Example
    - Database consisting of two items, X and Y
    - Only criterion for correctness:  $X = Y$
    - The following transactions (i.e., correct serial programs)
- |       |                      |       |                   |
|-------|----------------------|-------|-------------------|
| $T1:$ | $X \leftarrow X + 1$ | $T2:$ | $X \leftarrow 2X$ |
|       | $Y \leftarrow Y + 1$ |       | $Y \leftarrow 2Y$ |
- Initially,  $X=10$  and  $Y=10$ .

# Schedules

---

- When transactions execute concurrently, their operations are *interleaved*.
- A *schedule* is a sequence of operations from one or more transactions.
- Operations
  - *read*( $Q$ ,  $q$ )
    - ◆ Read the value of the database item  $Q$  and store in the local variable  $q$ .
  - *write*( $Q$ ,  $q$ )
    - ◆ Store in the database item  $Q$  the value of the local variable  $q$ .
  - Other operations such as arithmetic
  - *commit*
  - *rollback*

# Example

- One possible schedule:

- Initially,  
 $X=10$  and  $Y=10$ .

<i>Schedule S1</i>	
<i>T1</i>	<i>T2</i>
read(X,x) $x \leftarrow x+1$ write(X,x)	
	read(X,x) $x \leftarrow 2x$ write(X,x) read(Y,y) $y \leftarrow 2y$ write(Y,y)
read(Y,y) $y \leftarrow y+1$ write(Y,y)	

- Resulting database:  $X = 22$ ,  $Y = 21$ ,  $X \neq Y$ .

# The Concurrency Problem, cont.

---

- What do we mean by correctness?
- **Definition D1:** Concurrent execution of transactions must leave the database in a consistent state.
  - Assumes that each transaction, when started on a consistent state of the database leaves it in a consistent state.
- **Definition D2:** Concurrent execution of transactions must be (result) equivalent to some serial execution of the transactions.
  - *Result equivalent* means final database states must be identical.
- Which is better, **D1** or **D2**?

# Example

- Example, with an initial database satisfying  $X = Y$
- $S2$  is not result equivalent to a serial execution of  $T3$ ,  $T4$  even though the final database state is consistent.
  - Here only  $T4$  takes effect

Schedule $S2$	
$T3$	$T4$
read( $X, x$ )	
$x \leftarrow x+1$	
	read( $X, x$ )
write( $X, x$ )	
	$x \leftarrow 2x$
	write( $X, x$ )
	read( $Y, y$ )
	$y \leftarrow 2y$
read( $Y, y$ )	
$y \leftarrow y+1$	
write( $Y, y$ )	
	write( $Y, y$ )



# The Concurrency Problem, cont.

---

- Our choice is **D2**.
  - An execution sequence is *correct* if it is *result equivalent* to a serial execution.
- Suppose we are given a set of  $n$  transaction programs to be run concurrently *or* we can see a set of  $n$  transaction running concurrently. How do we check for correctness? This is not easy!
- A simplifying assumption:
  - We will look only at reads and writes on the databases to determine correctness.
  - This assumption is stronger than **D2**, as even fewer schedules are considered correct.

# Transactions

---

- The Transaction Concept
- The Concurrency Problem
- Serializability of Schedules
  - View Serializable
  - Conflict Serializable
- Recoverability

# View Equivalent Schedules

---

**Definition:** Two schedules  $S1$  and  $S2$  are *view equivalent* ( $S1 \equiv S2$ ) if

- The set of transactions participating in  $S1$  and  $S2$  are the same.
- For each data item  $Q$  in  $S1$ , if  $T_i$  reads  $Q$  and the value of  $Q$  read by  $T_i$  was written by  $T_j$ , then the same holds in  $S2$ .
  - ◆ This requirement ensures that the same values are read by all transactions in both  $S1$  and  $S2$ .
  - ◆ Therefore, the same computation occurs.
- For each data item  $Q$  in  $S1$ , if transaction  $T_i$  executes the last write of  $Q$ , then the same holds in  $S2$ .
  - ◆ This requirement ensures the same final system state by both schedules.

# View Serializability

- **Definition D3:** Let  $\{T_1, T_2, \dots, T_n\}$  be a set of transactions participating in schedule  $S$ .  $S$  is a *view serializable* schedule if there exists a serial schedule  $S_s$  such that  $S \equiv S_s$ .
- This provides a third notion of correctness (**D3**).
- A serializable (and serial) schedule:

Schedule S3		
T5	T6	T7
read(X,x) x ← x + 1 write(X,x)	read(X,x) x ← 2x write(X,x)	read(Y,y) y ← y + 1 write(Y,y)

# Example

- Is it a result serializable schedule (by **D2**)?
- Is it a view serializable schedule (by **D3**)?
- Recall that in D3, we only consider read and write operations, and can make no assumptions about the transactions' semantics.

Schedule S4	
T8	T9
read(X,x)	
$x \leftarrow x + 1$	
write(X,x)	
	read(Y,y)
	$y \leftarrow y - 10$
	write(Y,y)
read(Y,y)	
$y \leftarrow y + 1$	
write(Y,y)	
	read(X,x)
	$x \leftarrow x - 10$
	write(X,x)

# Example, cont.

---

- This schedule results in a correct database, but is not (view) serializable because of “lost information” in the definition.
- It is result serializable if we consider the values being stored by the final write operations (under **D2**). It yields the same results as  $T8, T9$  or  $T9, T8$ .
- But **D3** doesn't consider the values being stored.
- The operations we consider in a schedule w.r.t. **D3** are restricted to
  - read  $X$
  - write  $X$
  - commit
  - abort

# Possible Transaction Conflicts

- Assume that there are only two transactions,  $T1$  and  $T2$ , in the system.

$T1$	$T2$
write (X,x)	read (X,x)

$T1$	$T2$
read (X,x)	write (X,x)

$T1$	$T2$
write (X,x)	write (X,x)

$T1$	$T2$
read (X,x)	read (X,x)

# Conflict Equivalent Schedules

---

- Let  $I$  and  $J$  be consecutive instructions of a schedule  $S$  of different transactions.
- If  $I$  and  $J$  do not conflict, we can swap their order to produce a new schedule  $S'$ .
- The instructions appear in the same order in  $S$  and  $S'$ , except for  $I$  and  $J$ , whose order does not matter.
- $S$  and  $S'$  are termed *conflict equivalent schedules*.
- **Definition D4:** A schedule is *conflict serializable* if it is conflict equivalent to a serial schedule.



# Example

- $S5$  is view serializable, but is not conflict serializable, because every pair of consecutive instructions conflict.

view serializable?

conflict serializable?


Schedule $S5$		
$T1$	$T2$	$T3$
read(X,x)	write(X,x)	
write(X,x)		
		write(X,x)

- $T2$  and  $T3$  have write instructions without read instructions, termed *blind writes*.
- Blind writes appear in any view serializable schedule that is not conflict serializable.

# Conflict Graph

---

- We now proceed to construct a conflict (directed) graph for a schedule of a set of transactions.
  - We assume that a transaction will always read an item before it writes that item.
- Consider therefore some schedule of a set of transactions  $T1, T2, \dots, Tn$ .
  - The vertices of the conflict graph are the transaction identifiers.
  - An edge from  $Ti$  to  $Tj$  denotes the two transactions conflicting, with  $Ti$  making the relevant access earlier.
  - Sometimes the edge is labelled with the item involved in the conflict.
- Example for schedule  $S1$



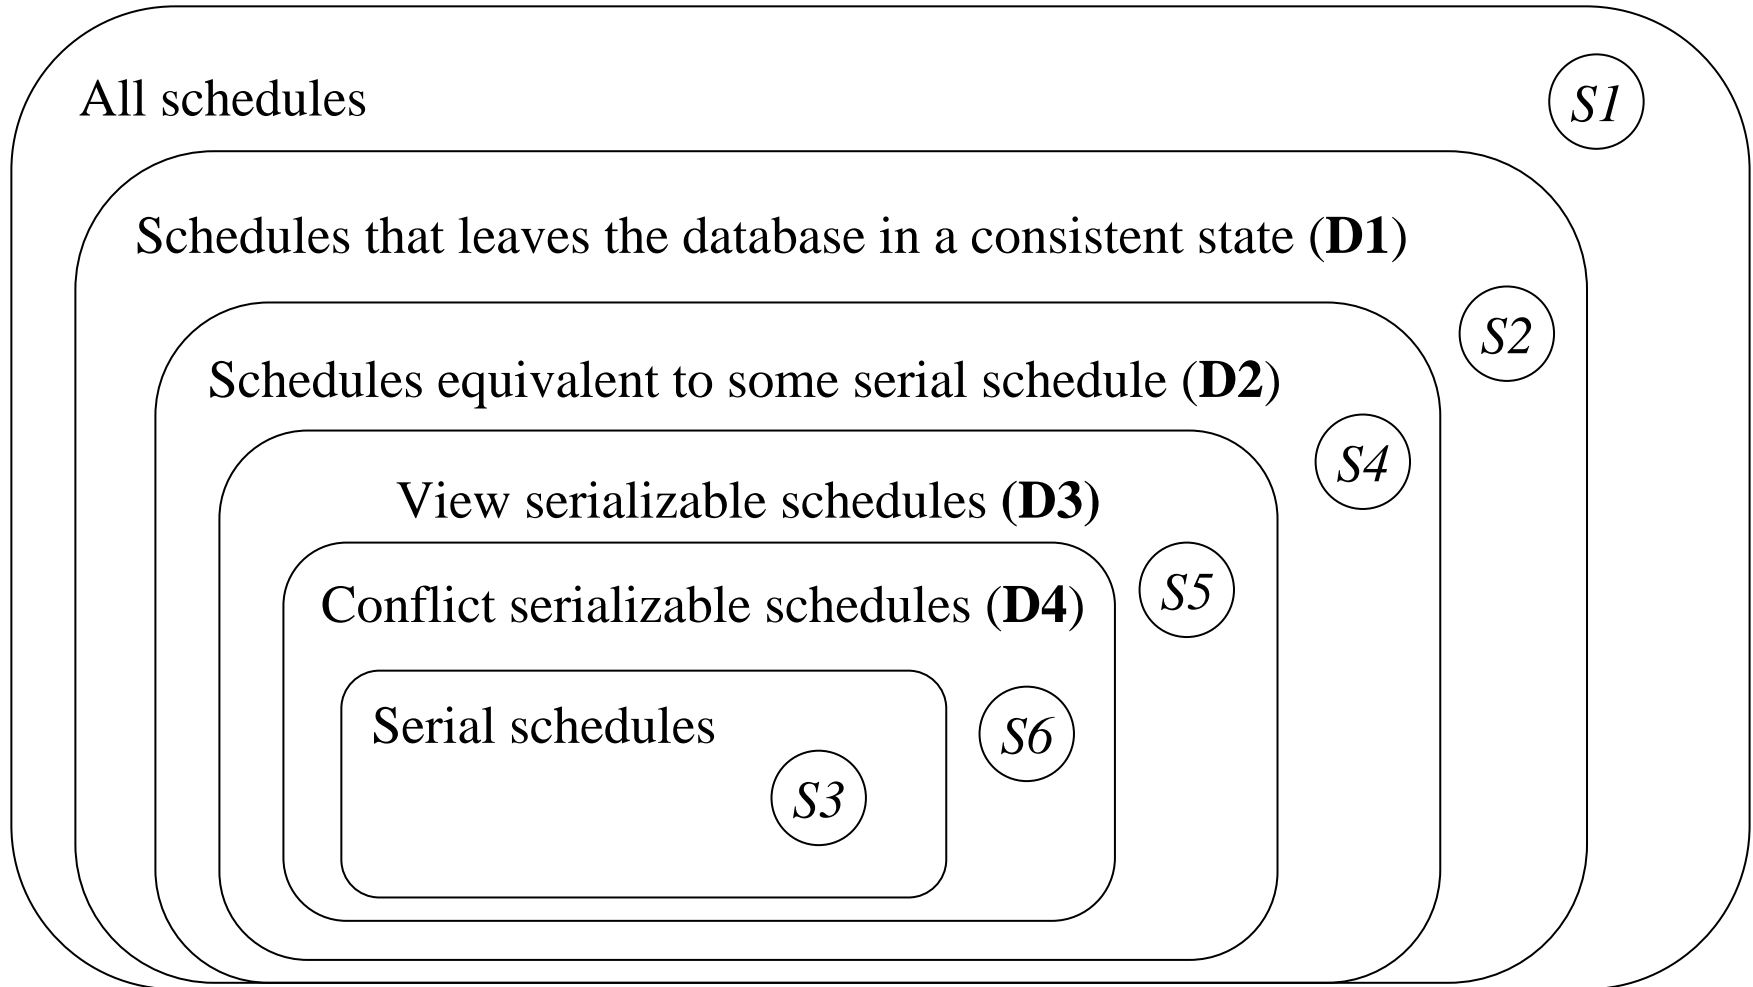
# Conflict Graph Example

Schedule S6				
T10	T11	T12	T13	T14
read Y read Z	read X			read V read W write W
	read Y write Y	read Z write Z		
read U			read Y write Y read Z write Z	
read U write U				

Exercises:

1. Draw out the conflict graph for this schedule.
2. Is it conflict serializable?

# Relationship Among Schedules



# Determining Serializability

---

- Given a schedule  $S$ , how do we determine that it is serializable?
- Use a slightly restricted definition: *conflict serializability* (**D4**).
- Two transactions  $T_i$  and  $T_j$  conflict if and only if there exists some item  $X$ , accessed by both  $T_i$  and  $T_j$ , and at least one of these transactions wrote  $X$ .
- Intuitively, a conflict between two transactions forces an execution order between them.
- We use conflict serializability, because it has a practical implementation.

# Testing Serializability

---

- A schedule is *conflict serializable* if its conflict graph is acyclic.
- Examples
  - Schedule *S1* had a graph containing a cycle, and is therefore not equivalent to any serial schedule.
  - Schedule *S5* had a graph which is acyclic, and is therefore equivalent to some serial schedule.
- Our goal: to develop protocols that will assure serializability (we will do this in the next lecture).
  - The protocols will generally not examine the conflict graph as it is being created.
  - Instead a protocol will impose the discipline that avoids non-serializable schedules.
- Testing for *view serializability* via conflict graphs is an NP-complete problem.

# Transactions

---

- The Transaction Concept
- The Concurrency Problem
- Serializability of Schedules
- Recoverability
  - Recoverable Schedules
  - Cascadeless Schedules

# Recoverable Schedules

---

- **Recoverable schedule:** For each pair of transactions  $T_i$  and  $T_j$ , where  $T_j$  reads data items written by  $T_i$ ,  $T_i$  must commit before  $T_j$  commits.
- The following schedule is not recoverable, if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read(A) write(A)  read(B)	  read(A)

- If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.



# Cascading Rollbacks

- If  $T_i$  fails it must be rollback to retain the atomicity property of transactions. If another transaction  $T_j$  has read a data item written by  $T_i$ ,  $T_j$  must also be rolledback.
- This is called *cascading rollback*, which we should avoid as it can lead to the undoing of a significant amount of work.
- Example
  - $T_{22}$  aborts.
  - We have to also rollback  $T_{23}$  and  $T_{24}$ .
  - They read “dirty” data.

Schedule S11		
$T_{22}$	$T_{23}$	$T_{24}$
read (A,a) read (B,b) write (A,a) write (B,b)	read (A,a)	read (A,a) read (B,b)
rollback		

# Cascadeless Schedules

- **Cascadeless schedule:** For each pair of transactions  $T_i$  and  $T_j$ , where  $T_j$  reads data items written by  $T_i$ , the commit operation of  $T_i$  must appear before the read operation of  $T_j$ .
  - Every cascadeless schedule is also recoverable
  - It is desirable to restrict the schedules to those that are cascadeless

Schedule S11		
T22	T23	T24
read (A,a)		
read (B,b)		
write (A,a)		
write (B,b)		
commit		
	read (A,a)	
		read (A,a)
		read (B,b)

# Summary

---

- Each transaction preserves database consistency.
- The serial execution of a set of transactions preserves database consistency.
- In a concurrent execution, steps of a set of transactions may be interleaved.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule.
- View serializability provides a convenient definition of correctness (NP-complete problem to check).
- Conflict serializability (has a practical implementation).
- Schedules must be cascadeless.